

# Partitioning From Zero to Hero



# In this session:



**Erik Darling Data**  
@erikdarlingdata



while brainstorming ideas for blog posts that talk about sql server consulting advice, i think the most common bit is also the shortest: you don't need partitioning.

4:22 pm · 17 Apr 2022 · Twitter Web App

16 Likes



# In this session:

- Some queries can be tuned to use a single partition (partition elimination)
- Can solve specific latch contention issues by spreading writes
- You can backup only read-write partitions – faster backups.
- You can restore the important partitions first – shorter RTO.



# Hi, I'm Daniel.

- SQL Server developer since 1999
- Owner and principal consultant of Structured Concepts AB
- Organizer of Data Saturday Stockholm
- Co-organizer of Group By
- Occasional blogger at [sqlsunday.com](http://sqlsunday.com)
- Author of [callfordataspeakers.com](http://callfordataspeakers.com)

Email: [daniel@strd.co](mailto:daniel@strd.co)

Twitter: [@dhmacher](https://twitter.com/dhmacher)

Blog: [sqlsunday.com](http://sqlsunday.com)



# Sponsors rock our world.



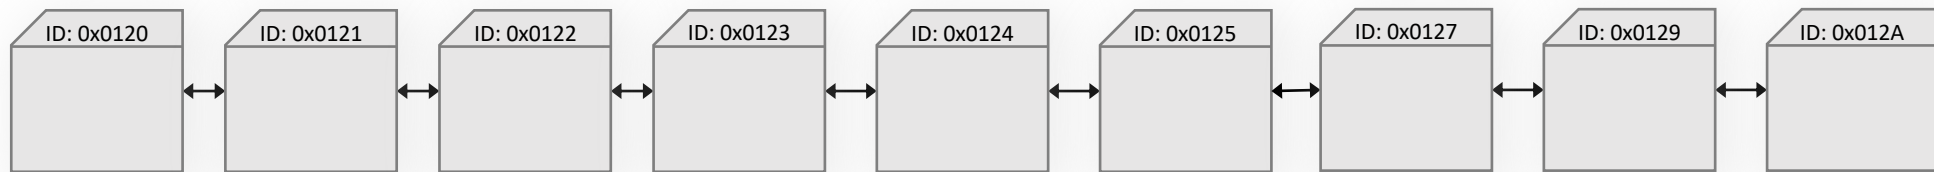
Denny Cherry  
& Associates Consulting



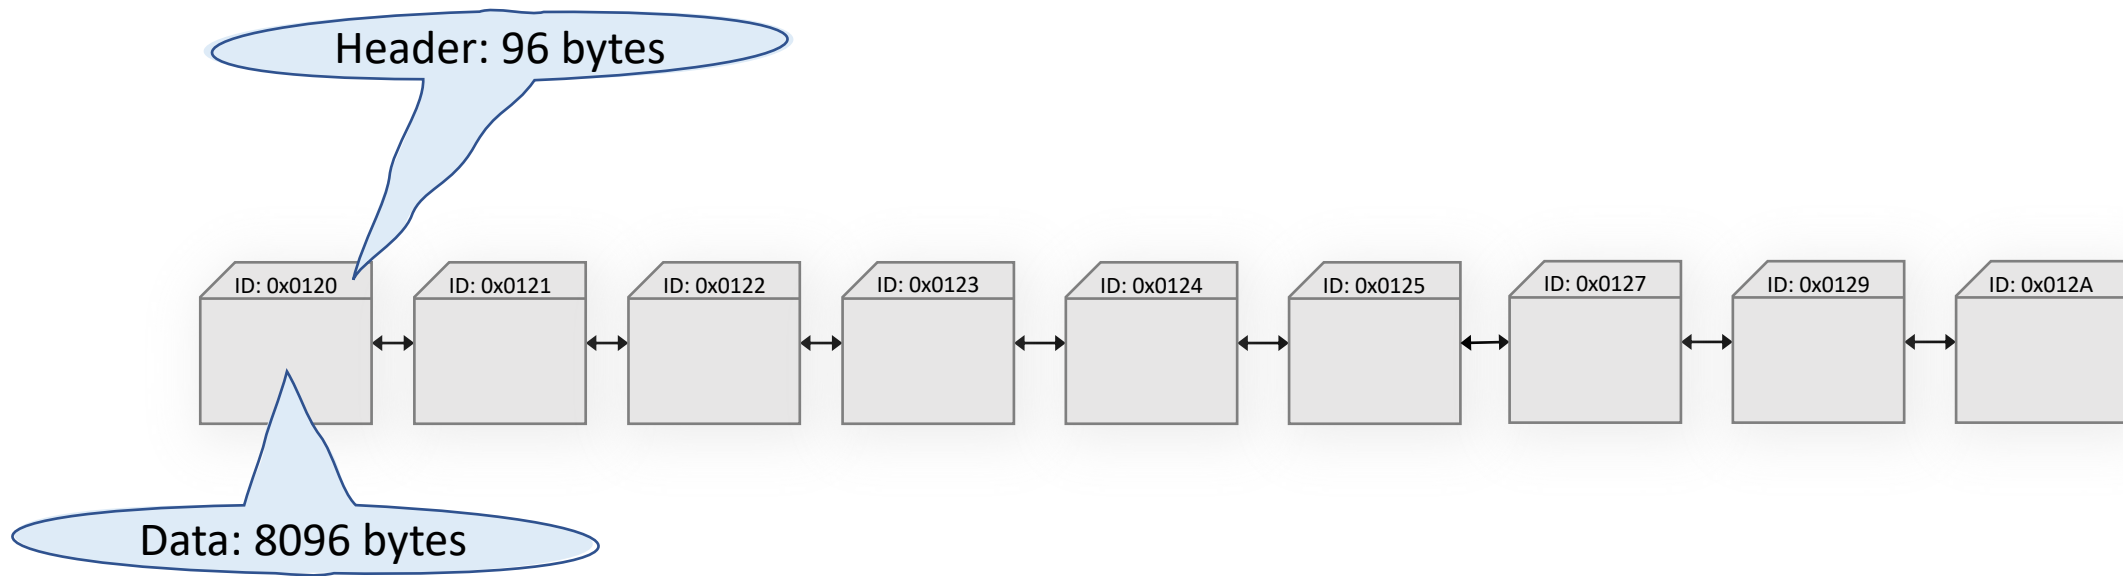
Simplifications ahead



# How pages work

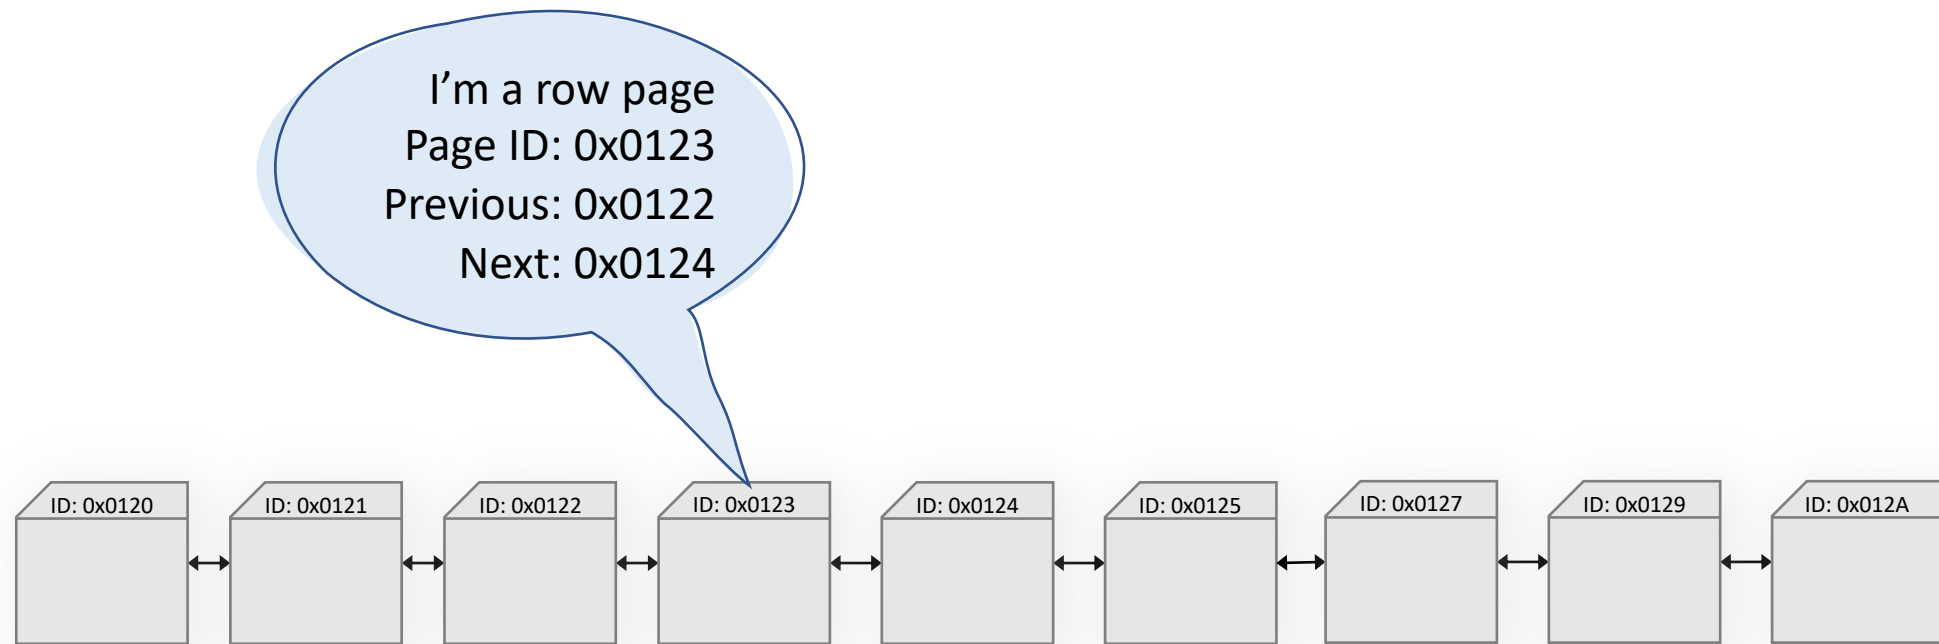


# How pages work

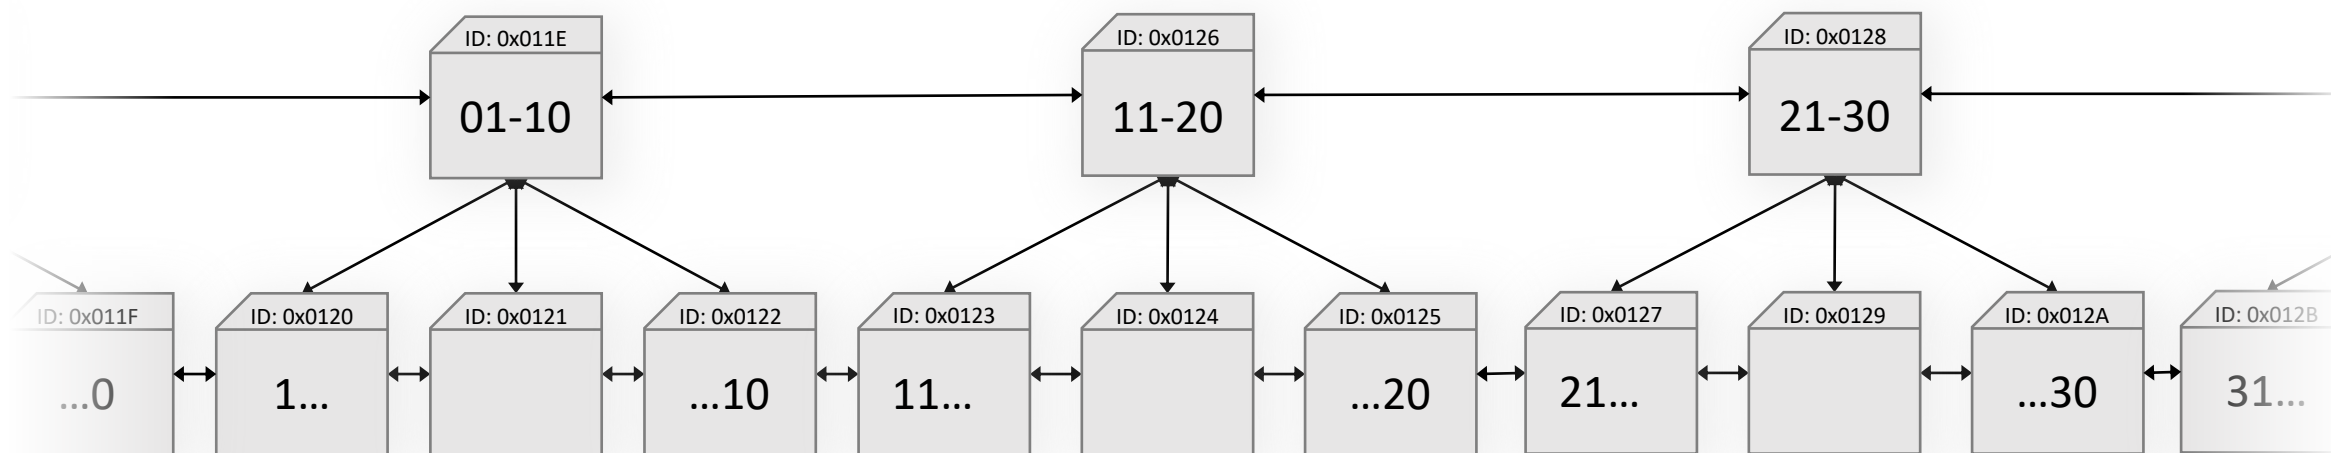




# How pages work

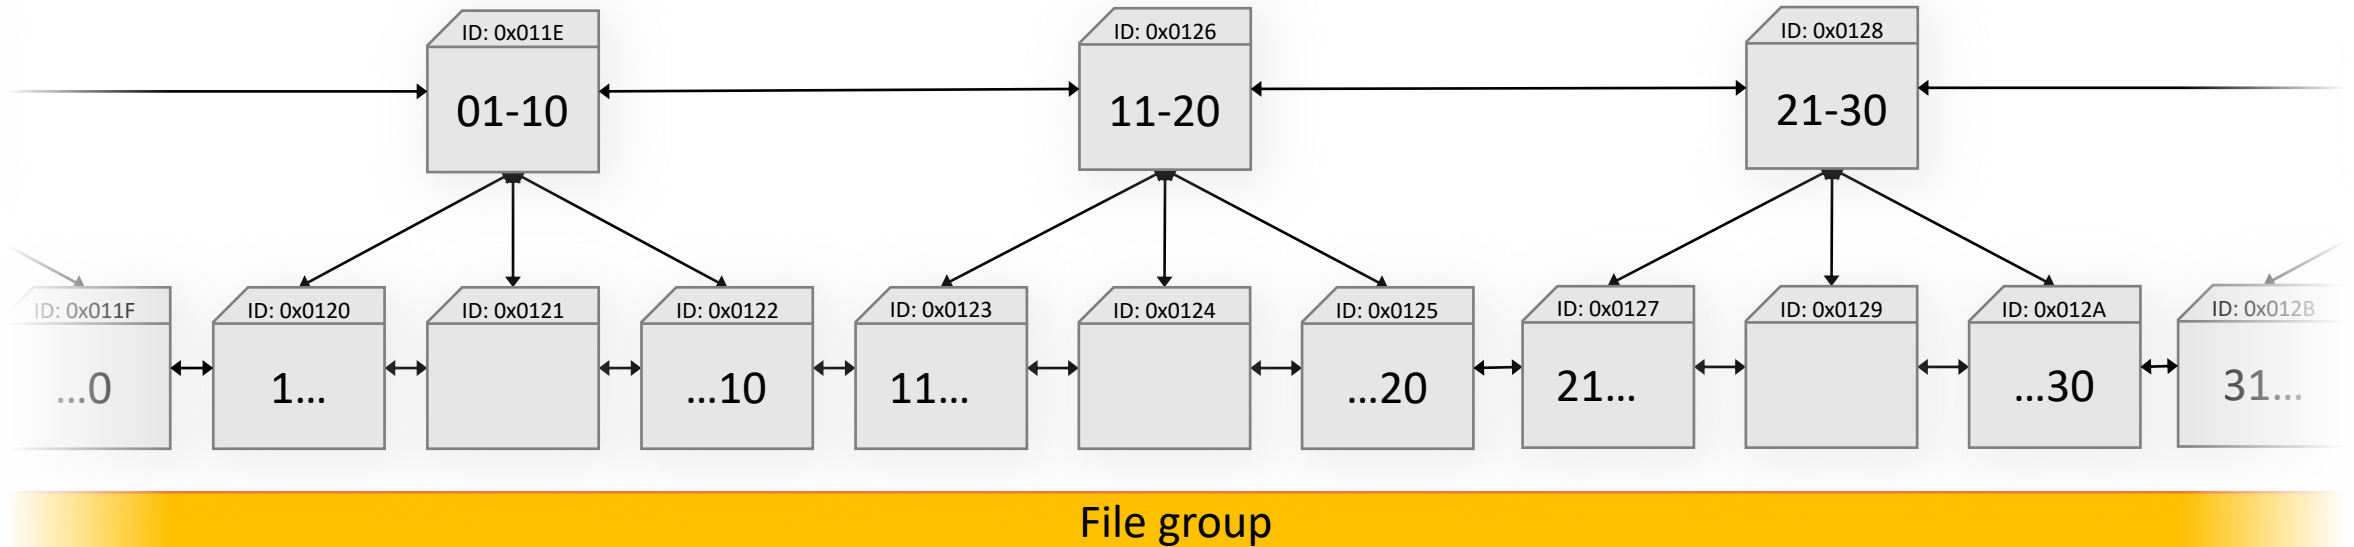


# Indexes can have multiple levels



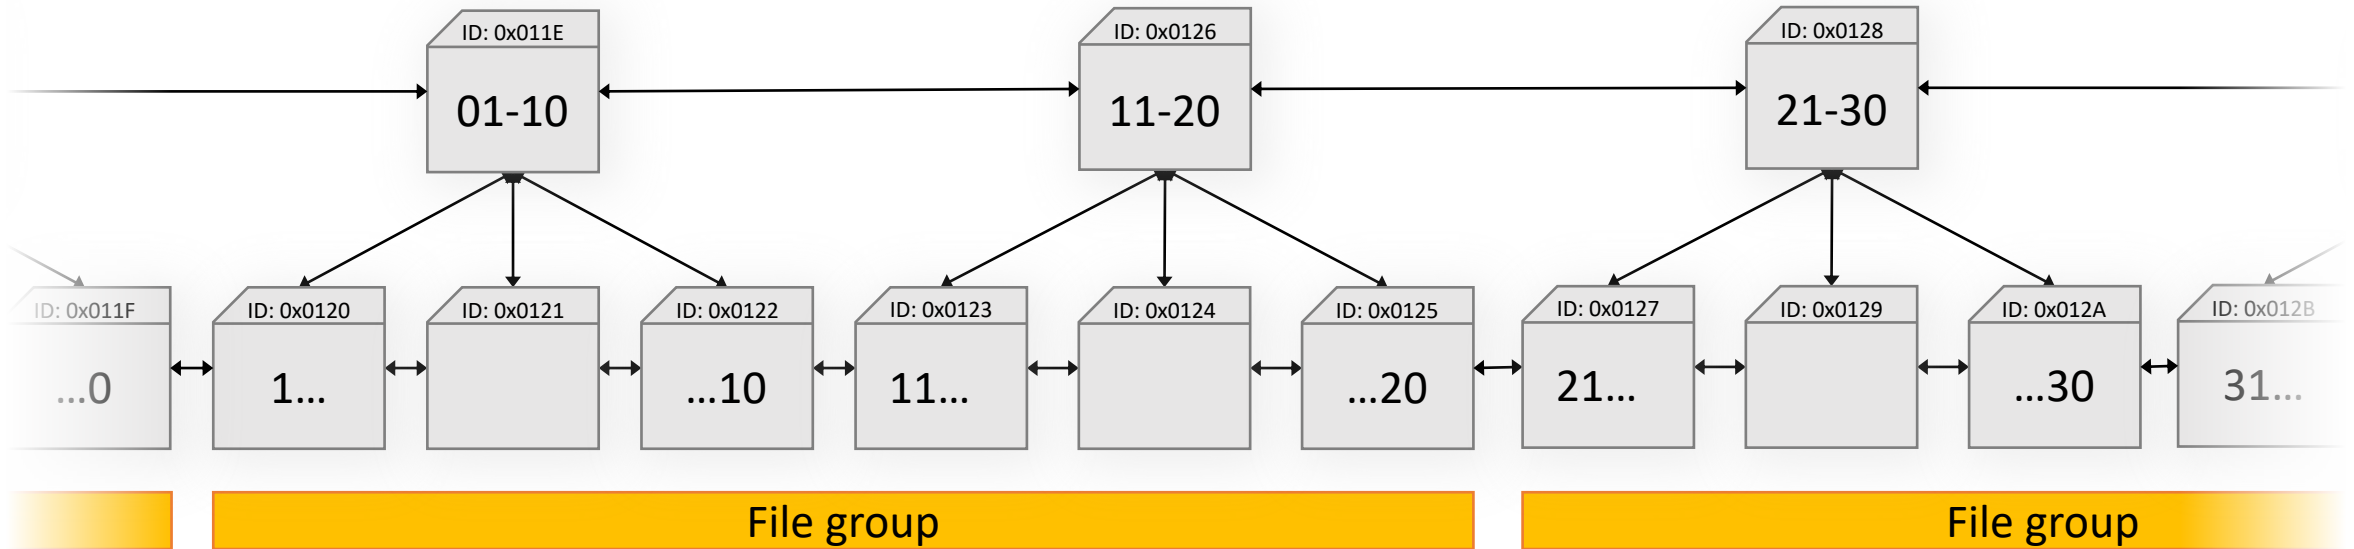
# What is partitioning?

Regular table: one partition on one filegroup



# What is partitioning?

Partitioned table: *multiple* partitions, on *one or more* filegroups.



# Setting up partitioning



*How: the partition function*

```
CREATE PARTITION FUNCTION AnnualFunction(date)
AS RANGE RIGHT
FOR VALUES ( '2020-01-01' , '2021-01-01' , '2022-01-01' );
```



# Setting up partitioning



*How: the partition function*

```
CREATE PARTITION FUNCTION AnnualFunction(date)
AS RANGE RIGHT
FOR VALUES ( '2020-01-01' , '2021-01-01' , '2022-01-01' );
```

2020-01-01 – 2020-12-31

2021-01-01 – 2021-12-31

2022-01-01 – 2022-12-31

# Setting up partitioning

```
CREATE PARTITION FUNCTION AnnualFunction(date)
AS RANGE RIGHT
FOR VALUES ( '2020-01-01' , '2021-01-01' , '2022-01-01' );
```

```
CREATE PARTITION SCHEME Annual
AS PARTITION AnnualFunction
TO ( [FILEGROUP_A] , [FILEGROUP_B] , [FILEGROUP_C] , [FILEGROUP_D] );
```

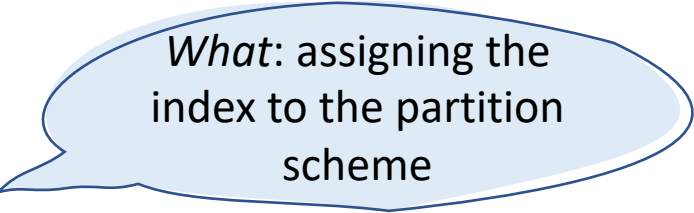
Where: the partition  
scheme

# Setting up partitioning

```
CREATE PARTITION FUNCTION AnnualFunction(date)
AS RANGE RIGHT
    FOR VALUES ( '2020-01-01' , '2021-01-01' , '2022-01-01' );
```

```
CREATE PARTITION SCHEME Annual
AS PARTITION AnnualFunction
TO ( [FILEGROUP_A] , [FILEGROUP_B] , [FILEGROUP_C] , [FILEGROUP_D] );
```

```
CREATE TABLE dbo.partitioned_table (
    TransactionDate    date NOT NULL,
    ...
    ...
    ...
) ON Annual(TransactionDate);
```



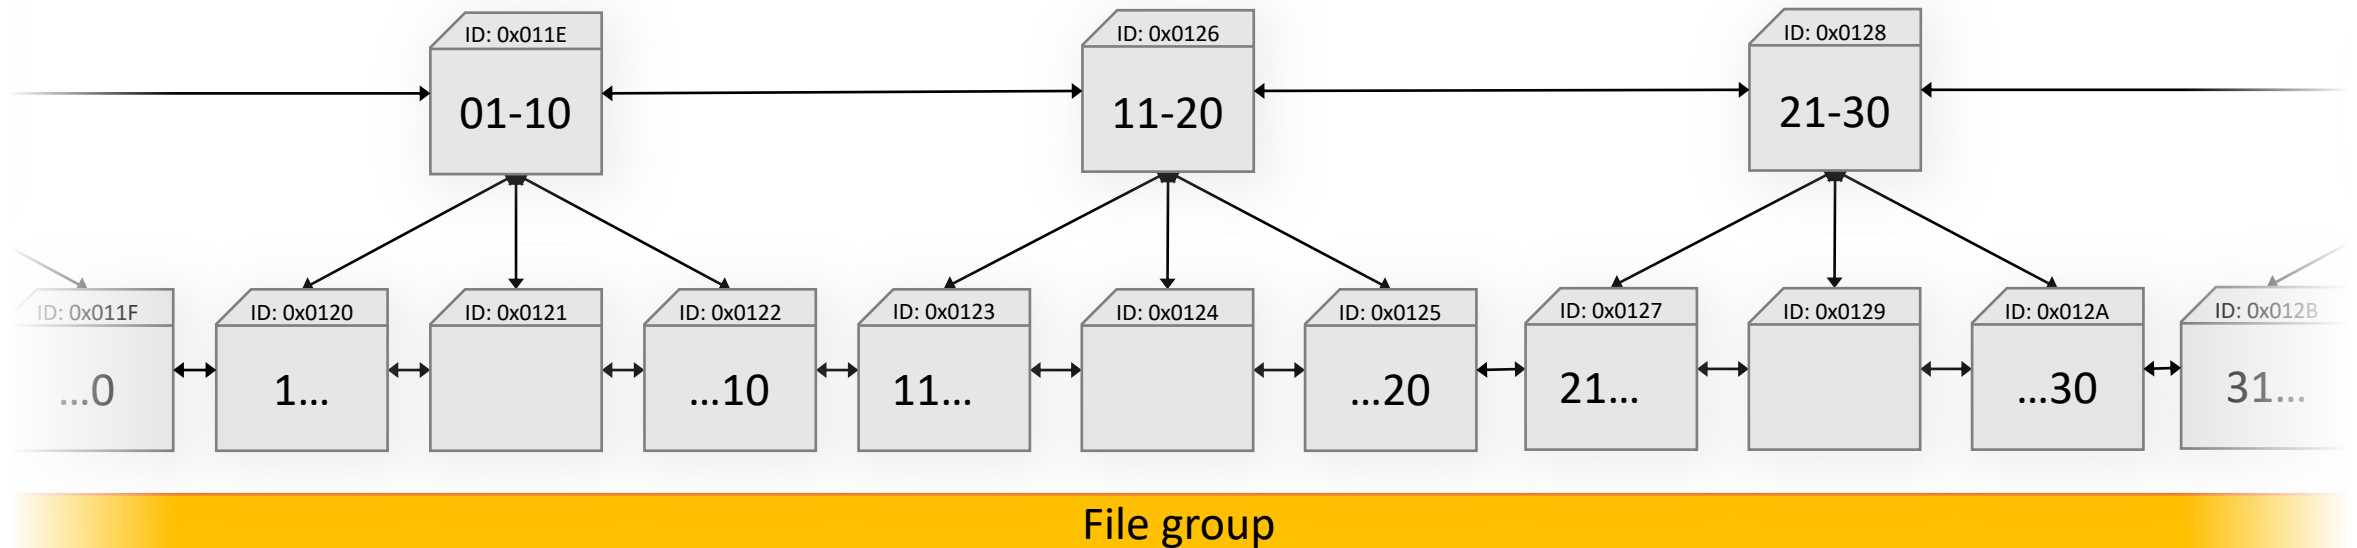
What: assigning the index to the partition scheme



DEMO

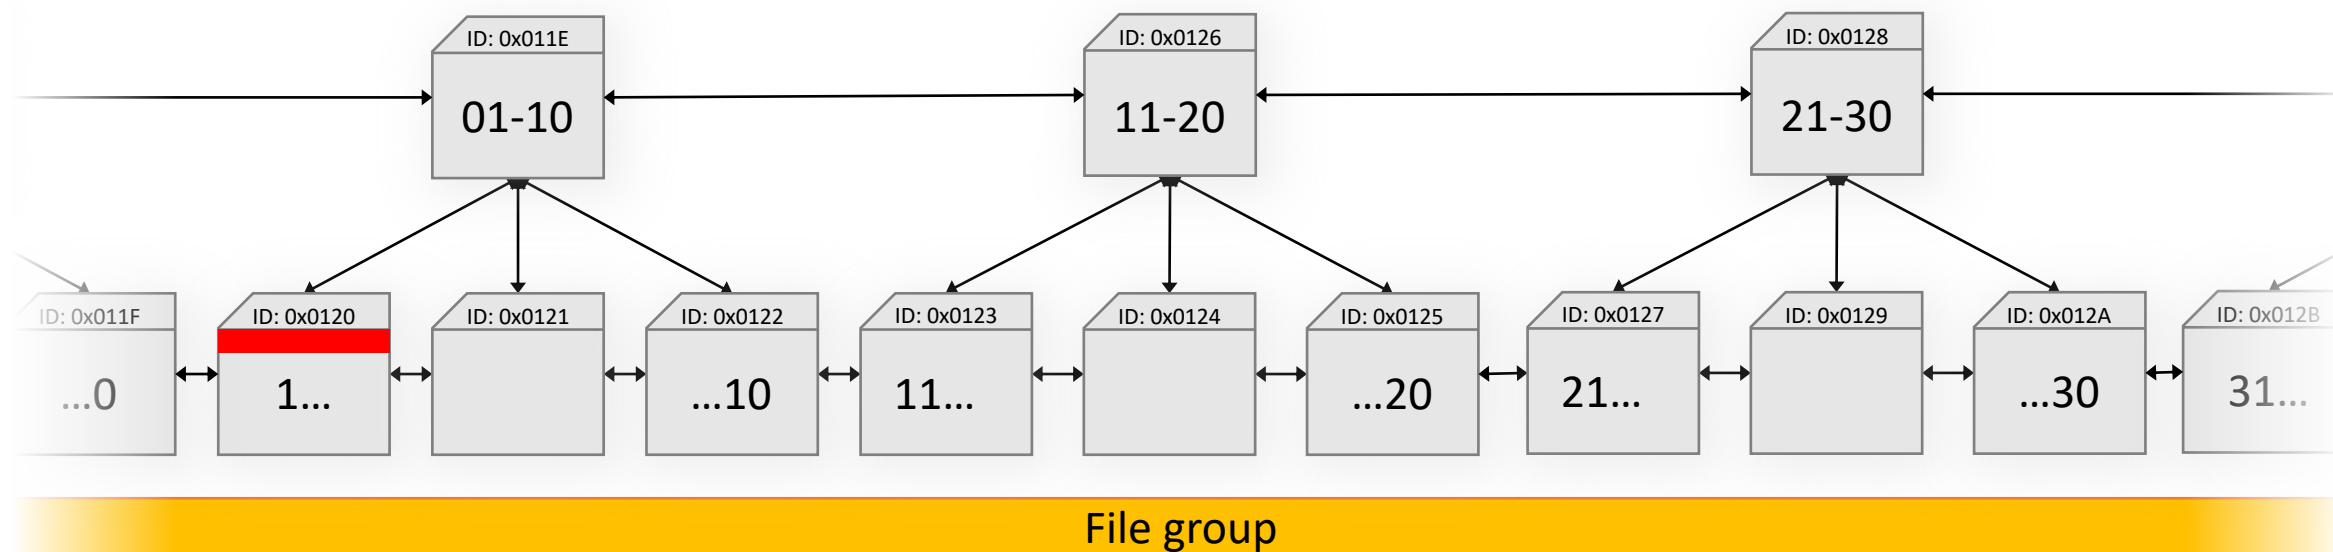
# A regular delete

```
DELETE FROM dbo.regular_table  
WHERE id BETWEEN 1 AND 20;
```

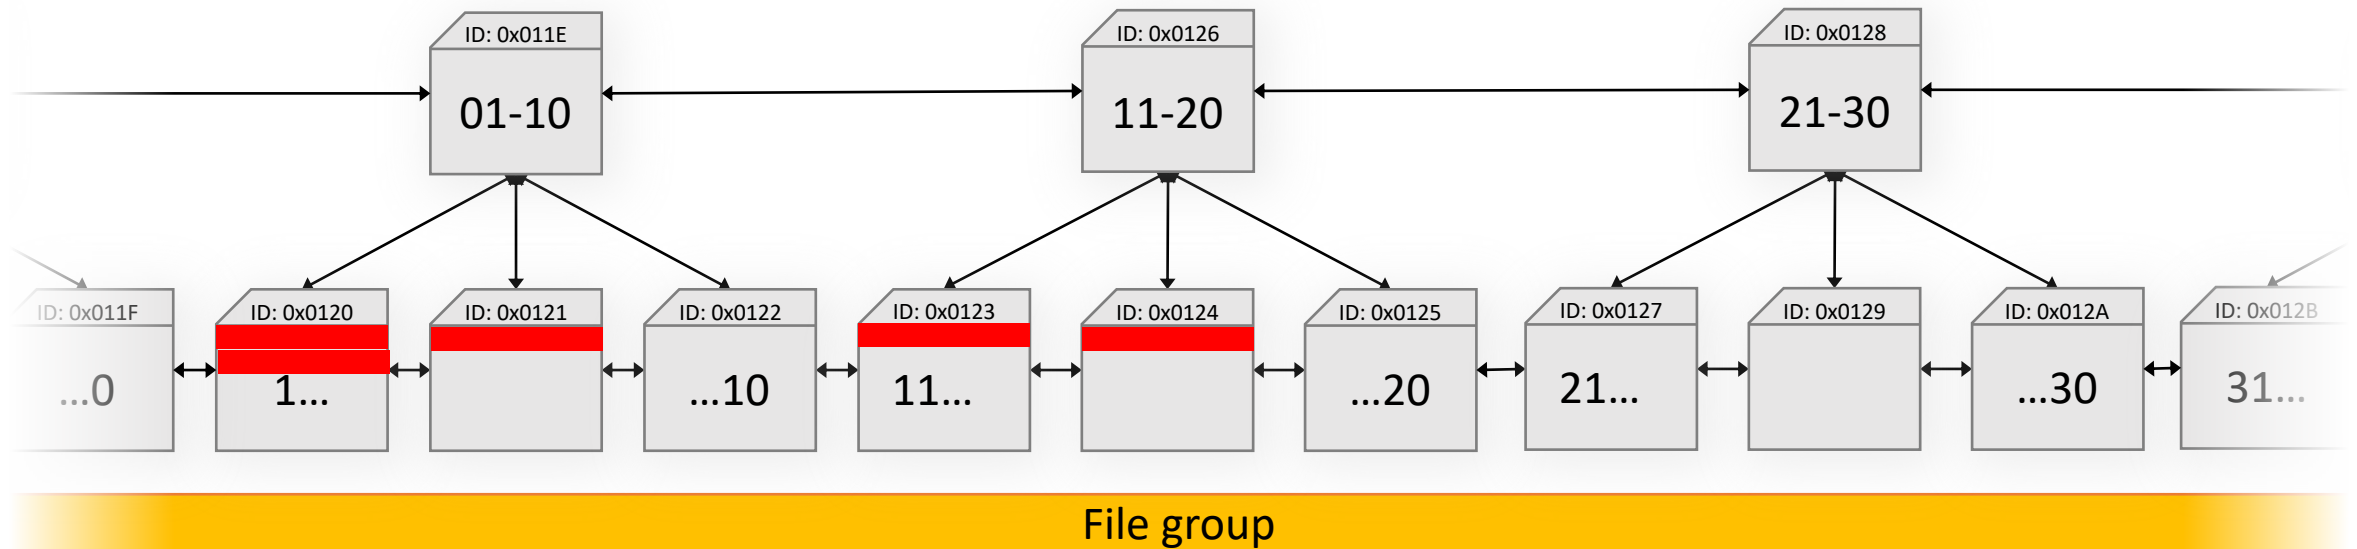




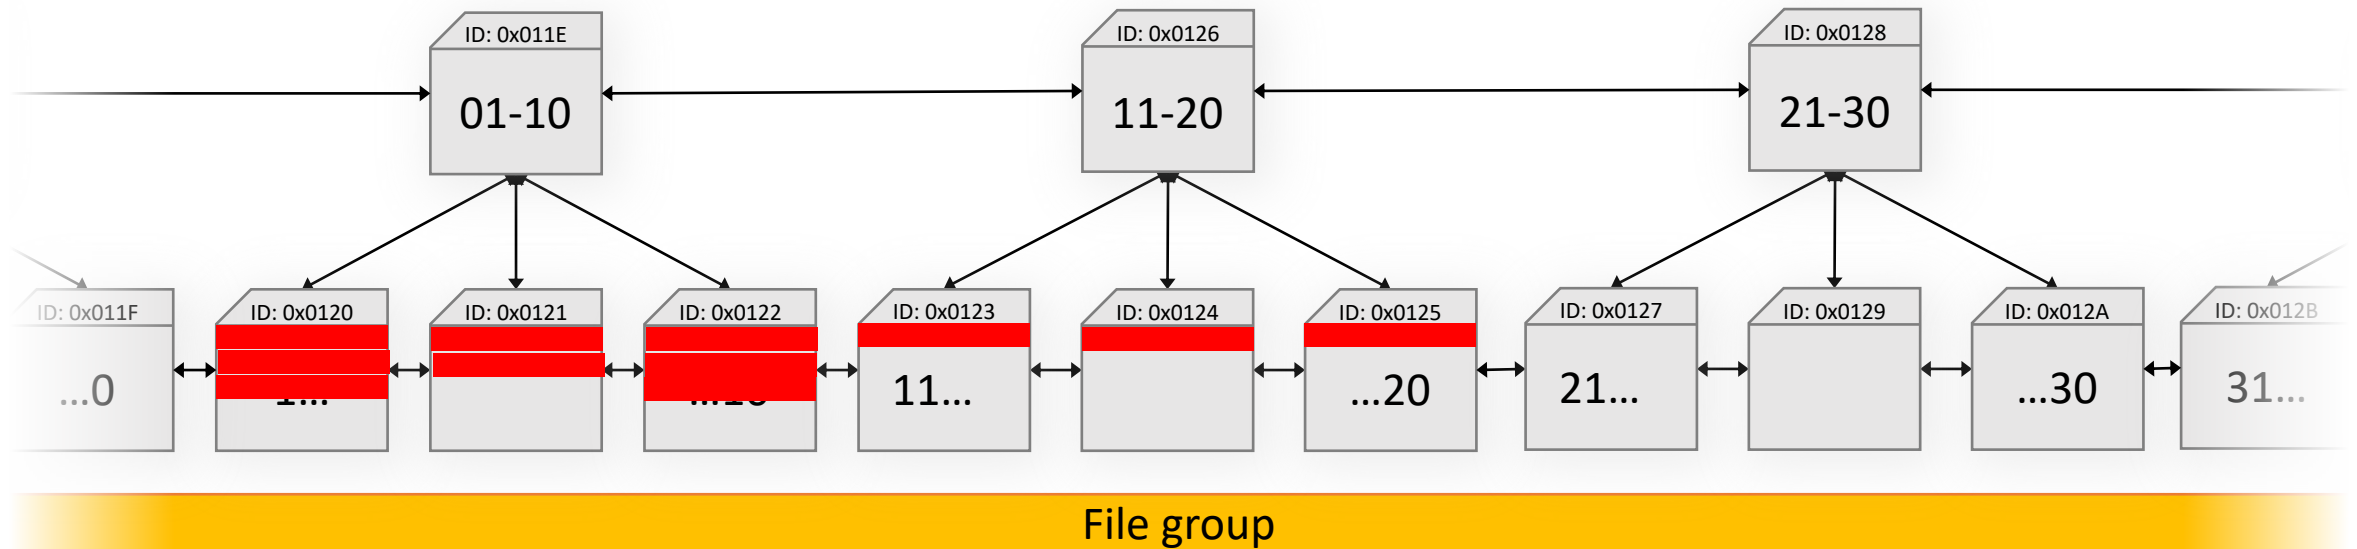
# A regular delete



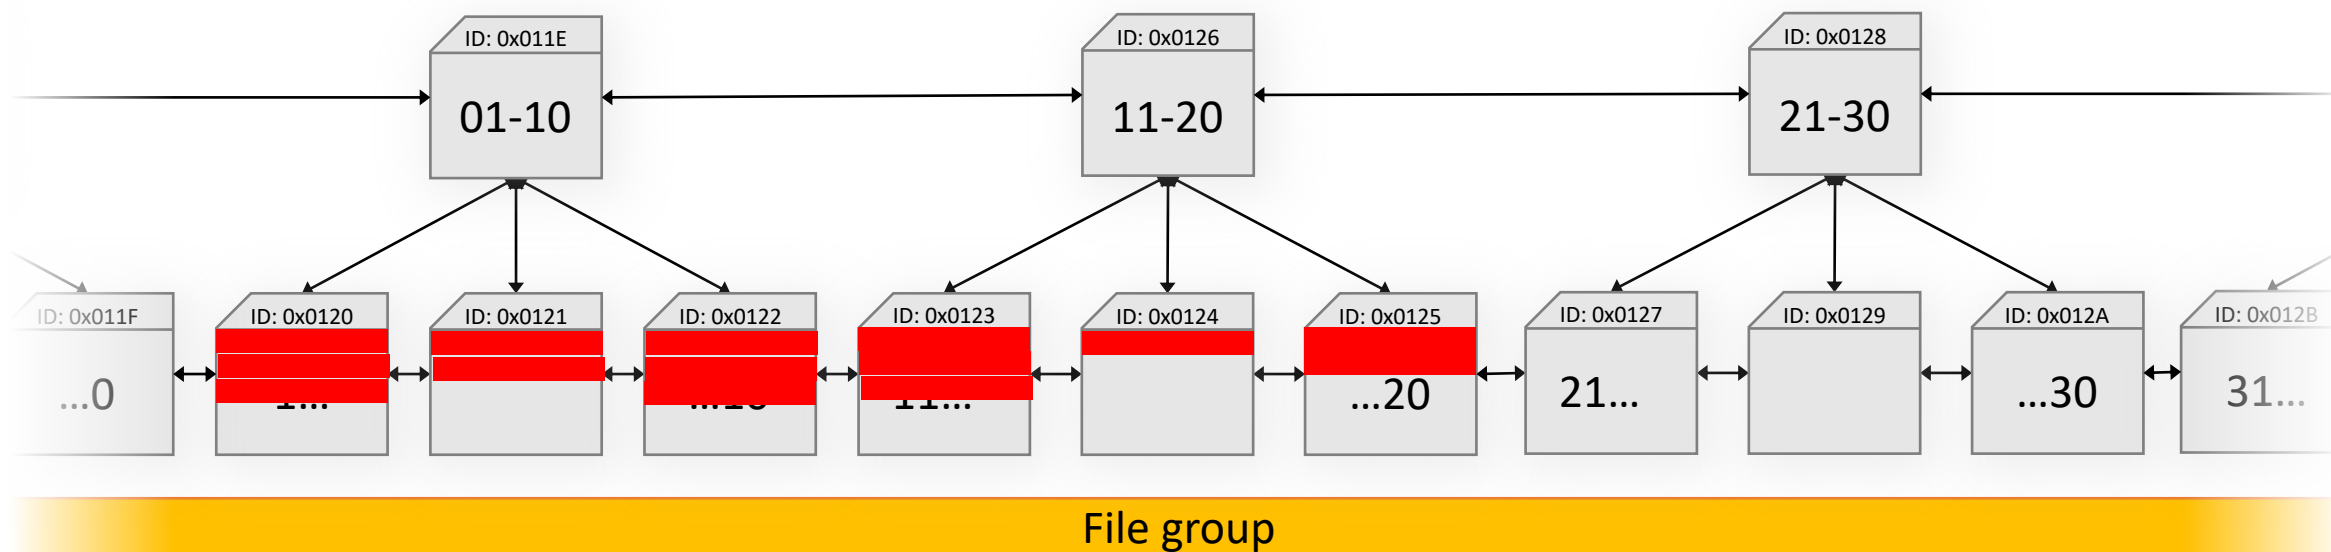
# A regular delete



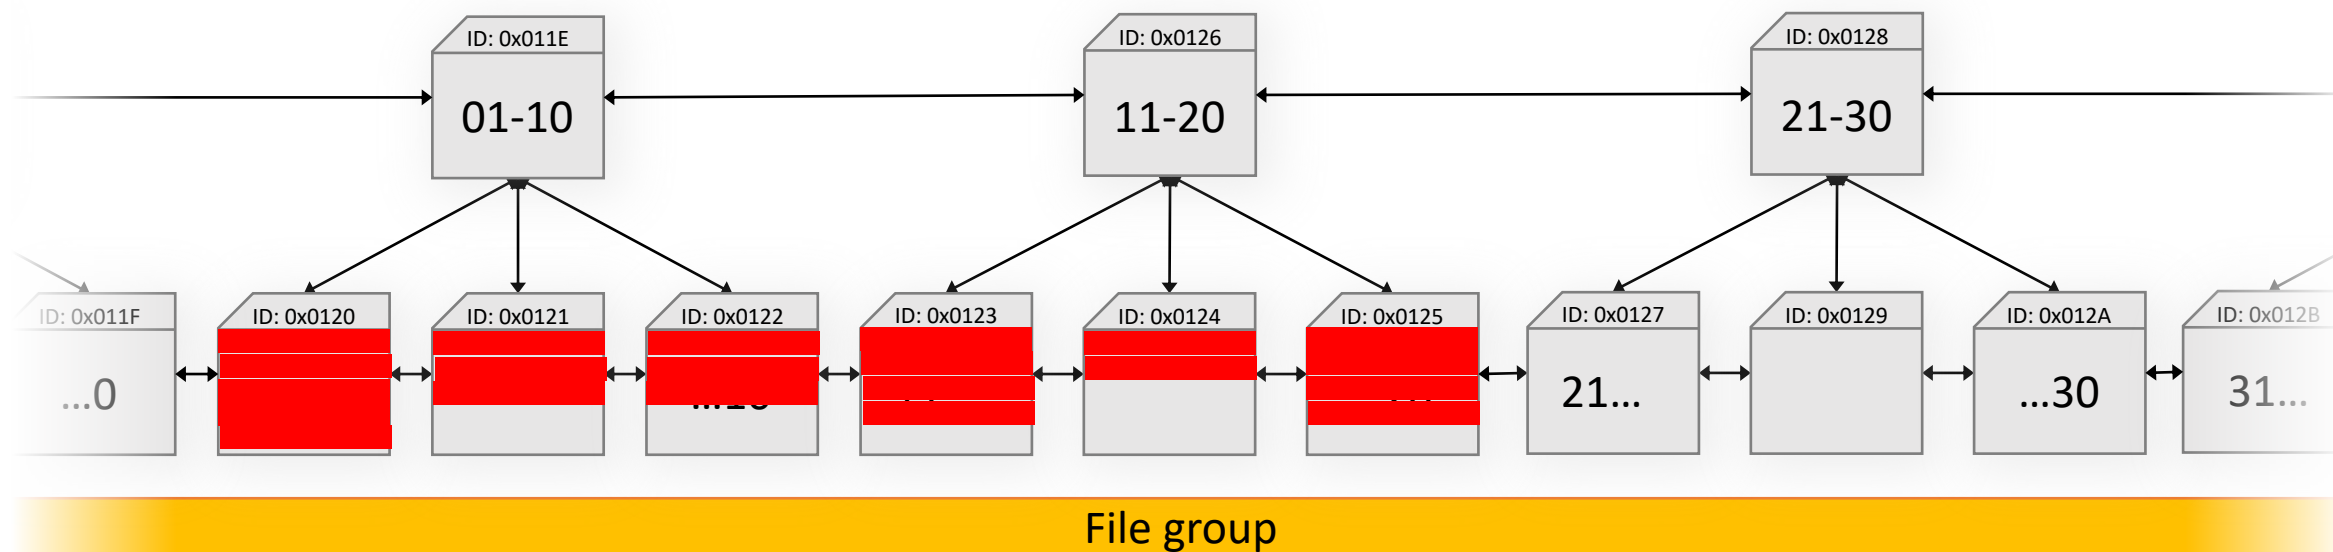
# A regular delete



# A regular delete

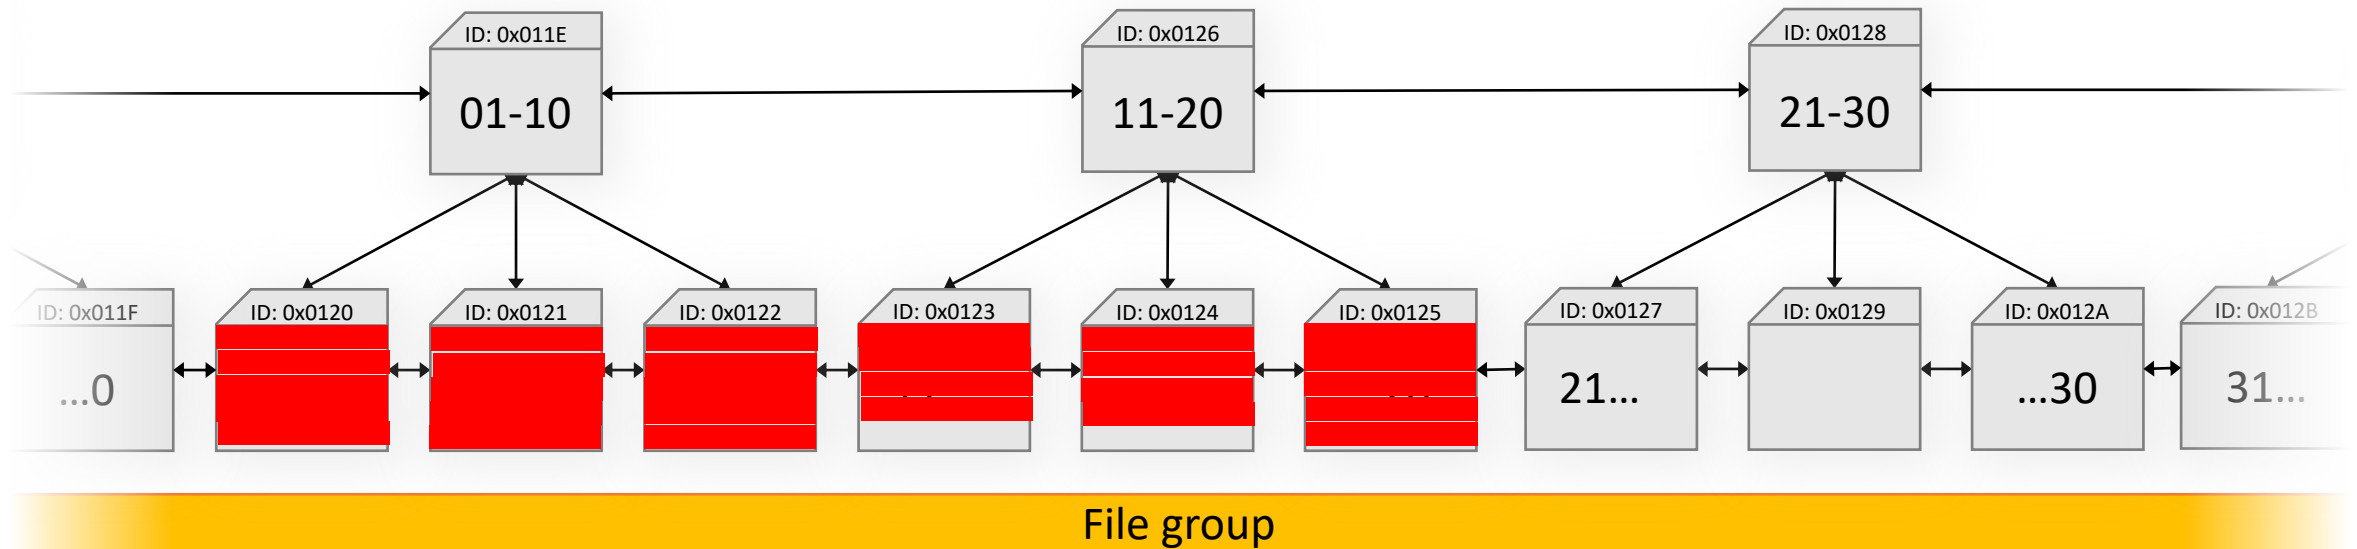


# A regular delete

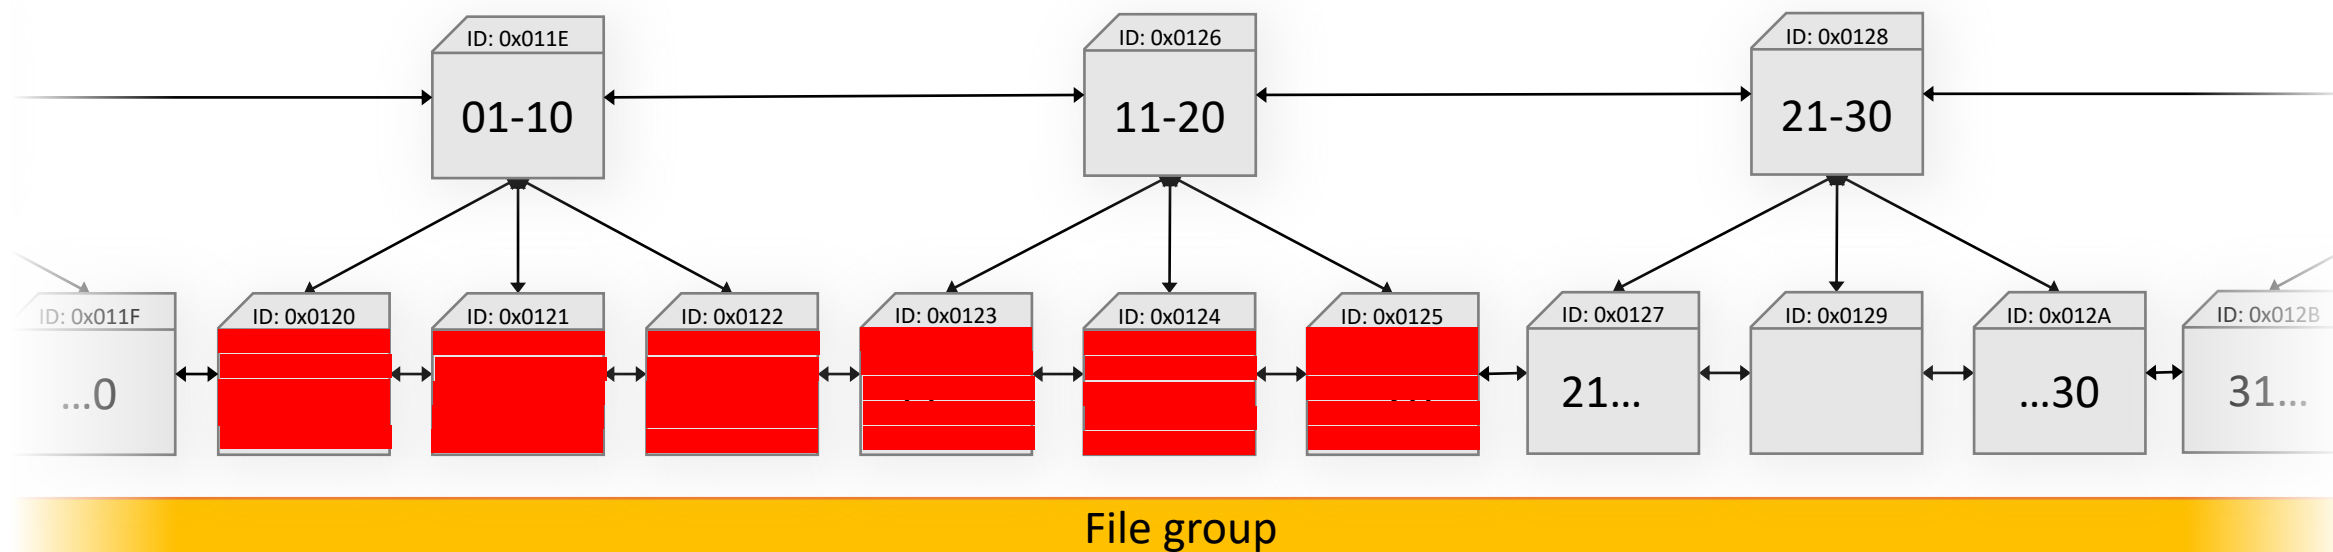




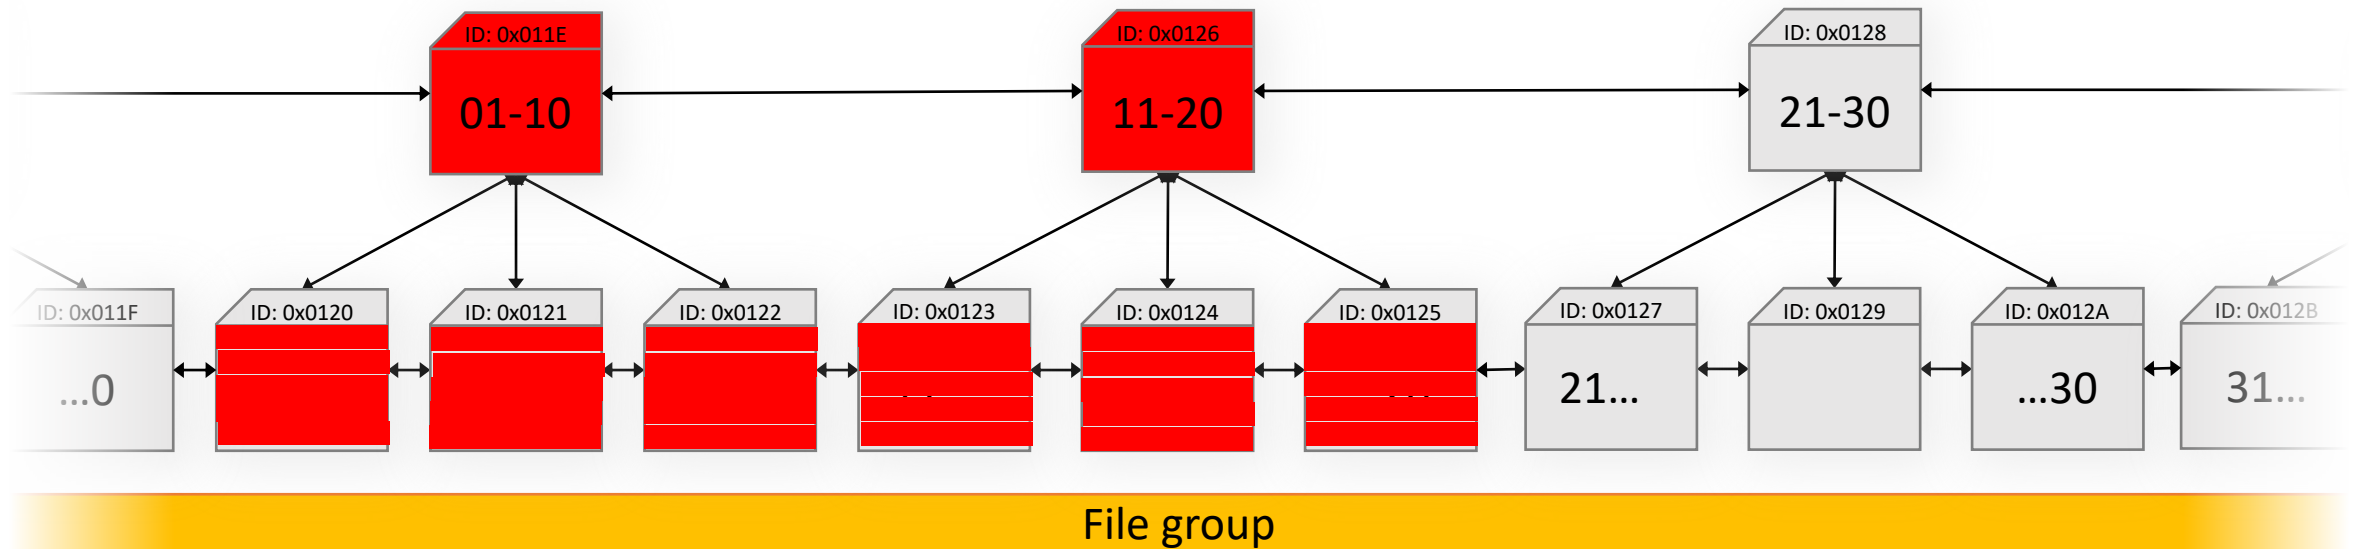
# A regular delete



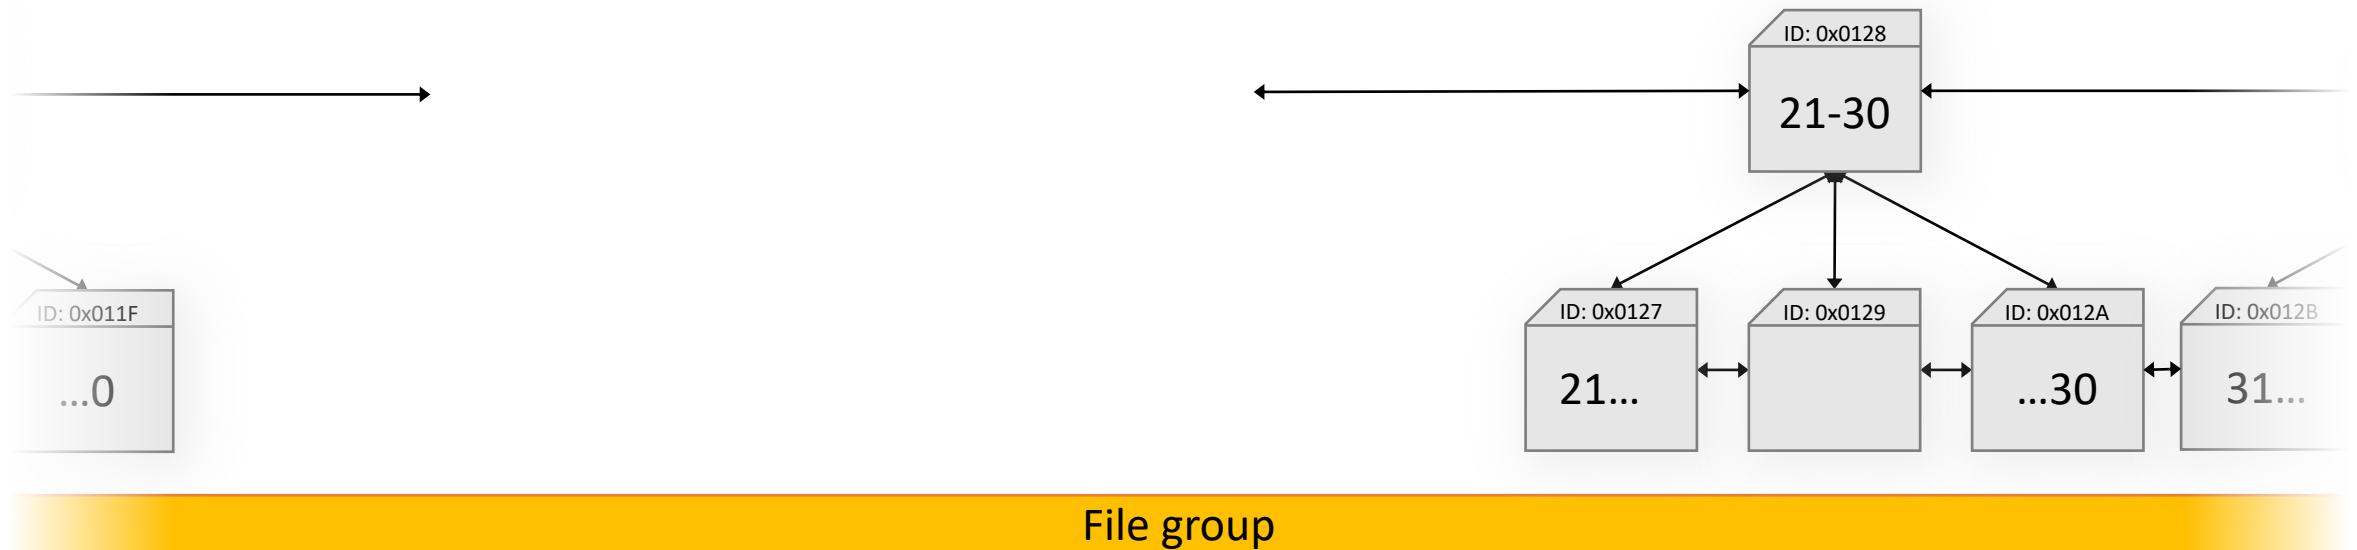
# A regular delete



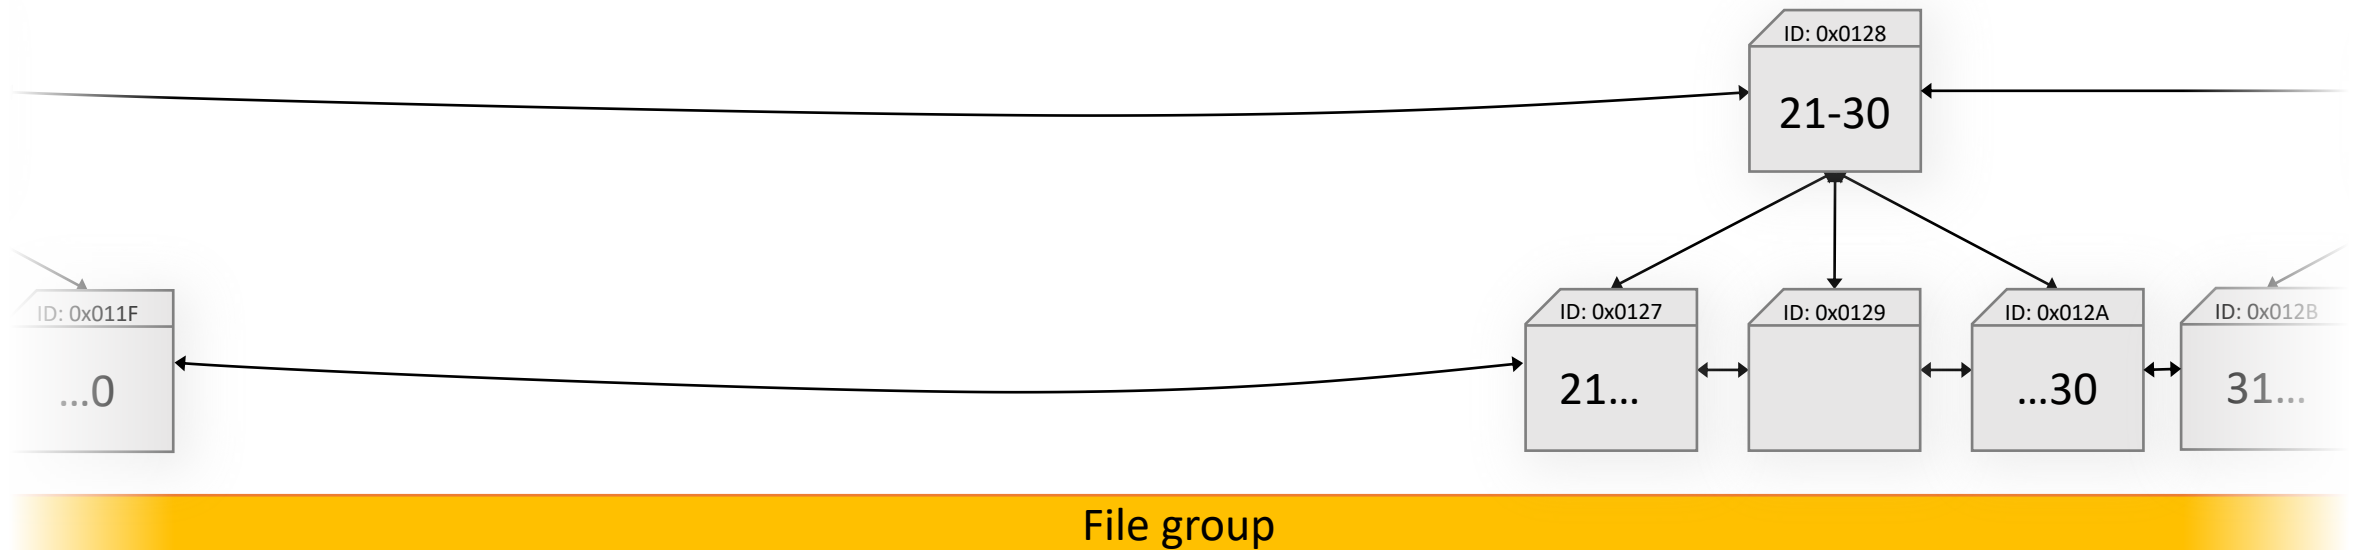
# A regular delete



# A regular delete



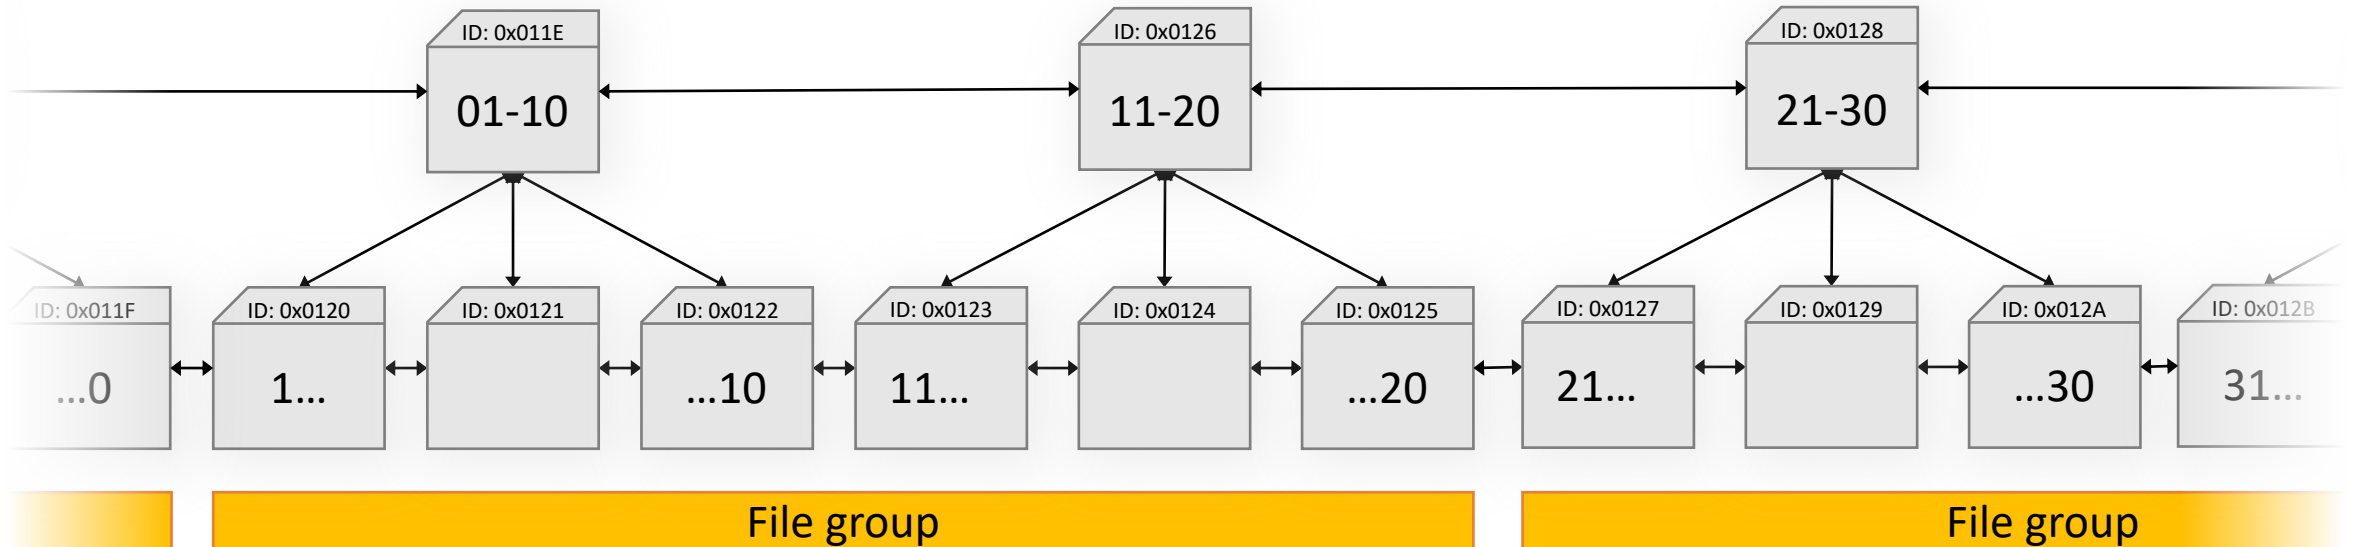
# A regular delete





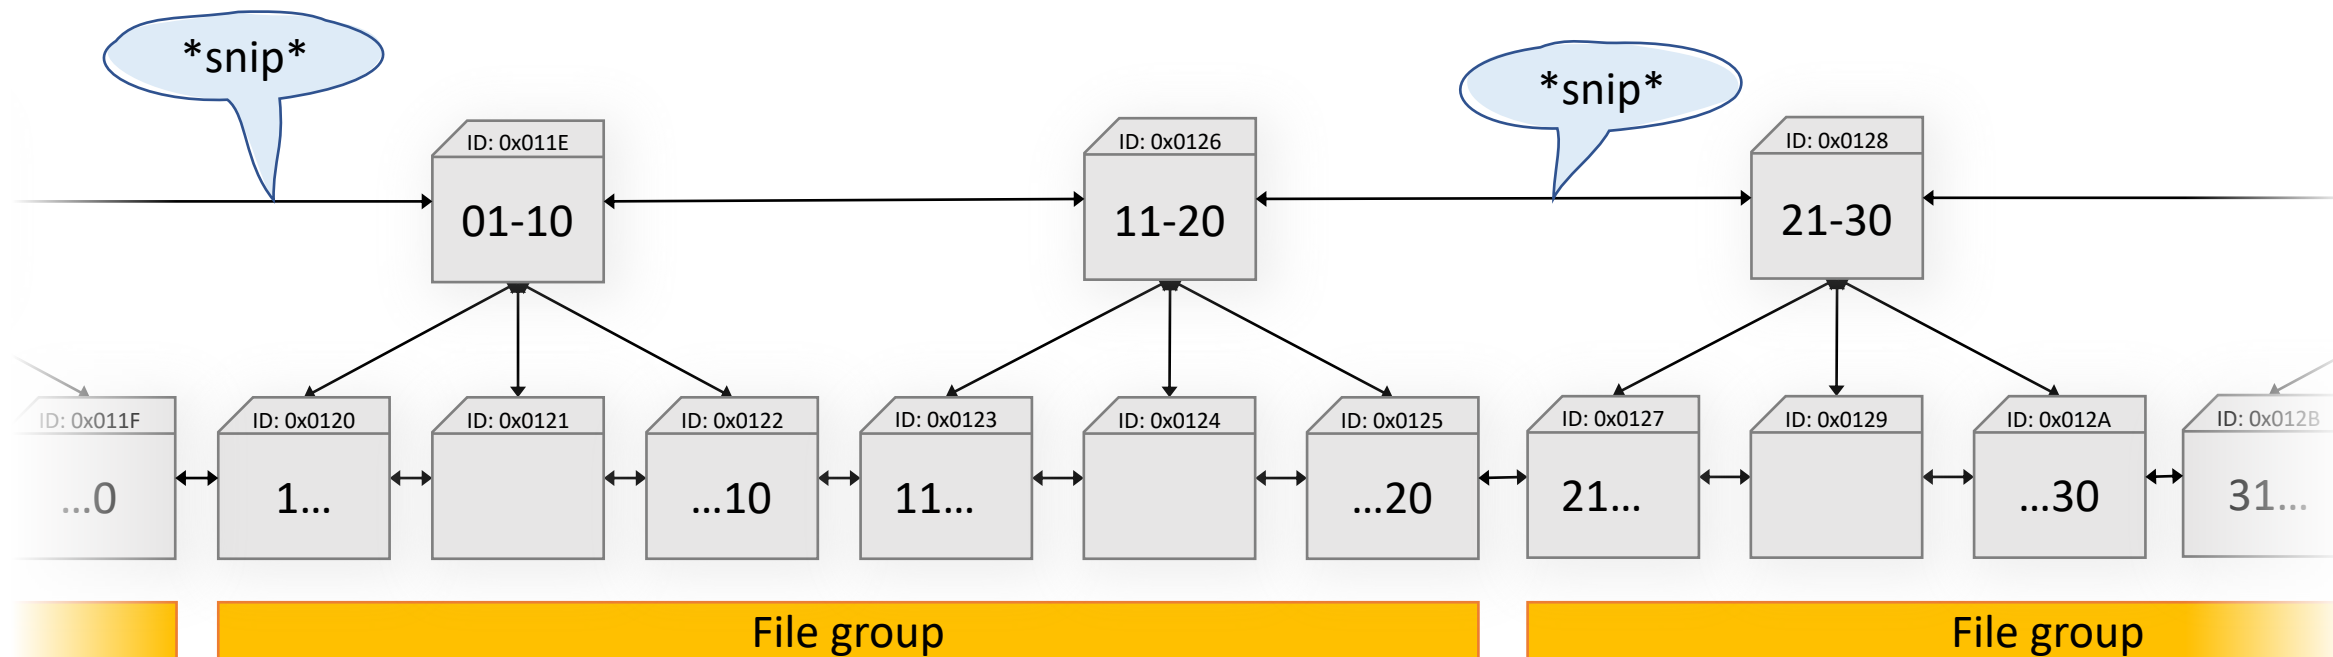
# Truncating a partition

```
TRUNCATE TABLE dbo.partitioned_table  
WITH (PARTITIONS (2));
```

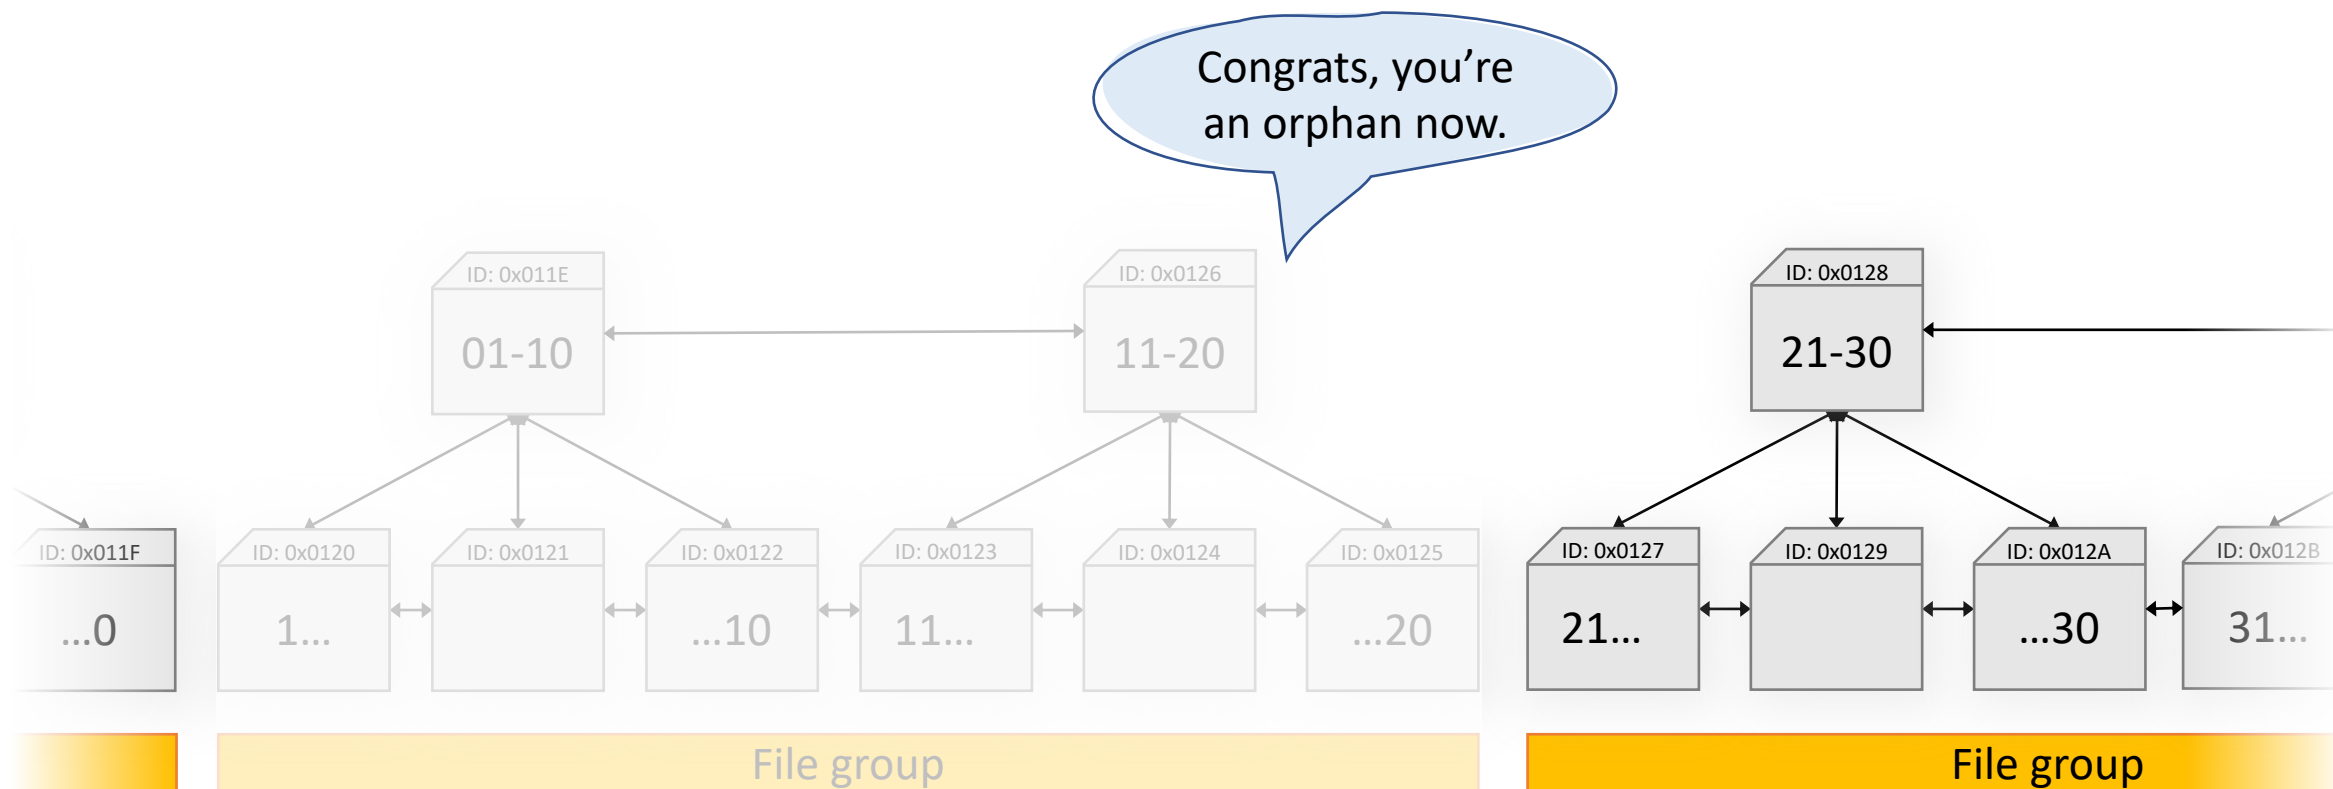


\* Truncating individual partitions is available on SQL Server 2016+

# Truncating a partition

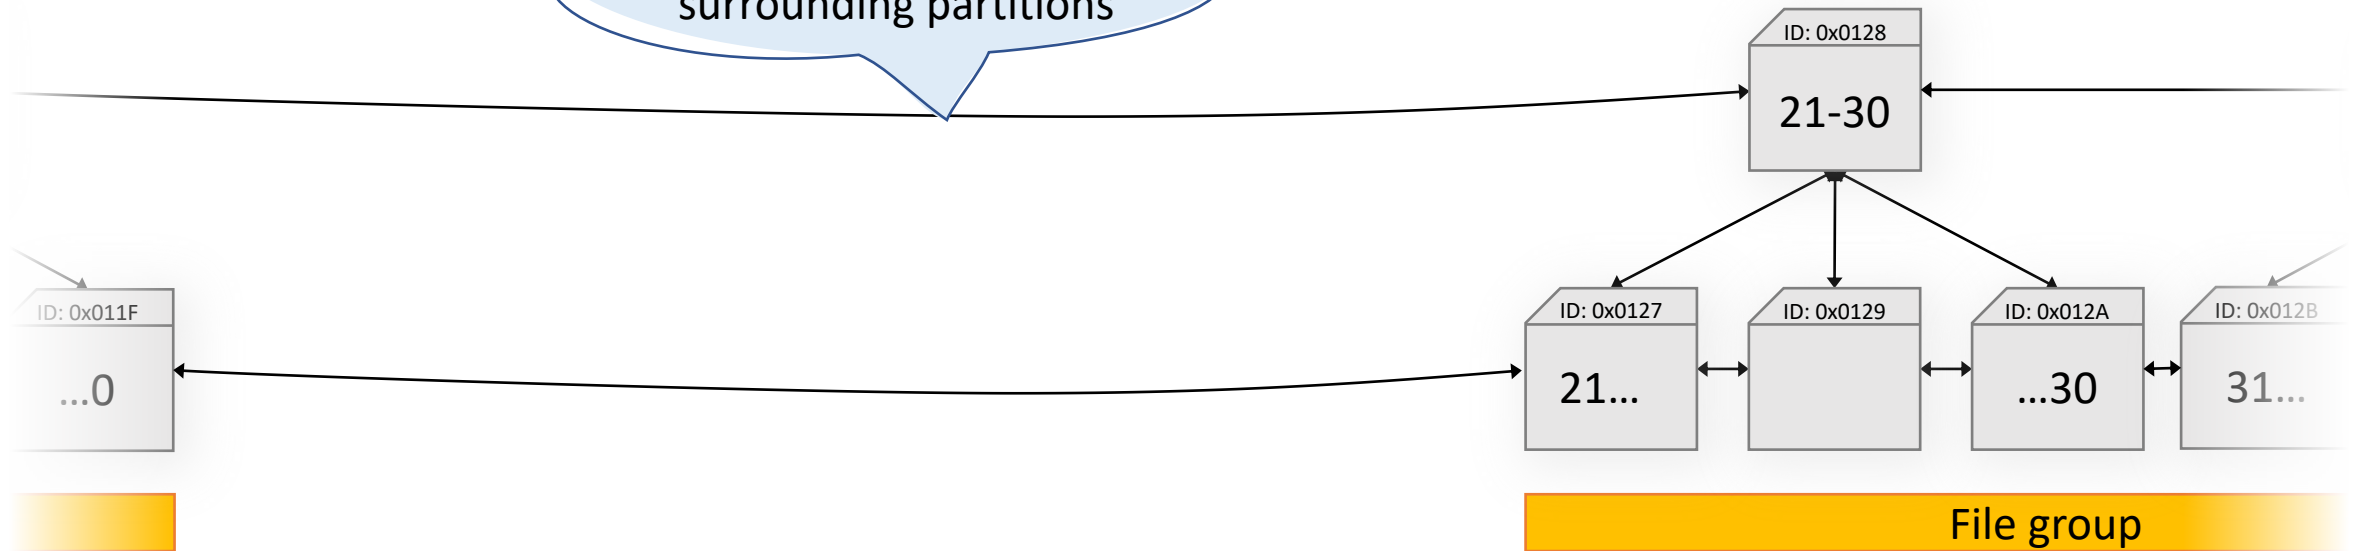


# Truncating a partition



# Truncating a partition

Connecting the  
surrounding partitions

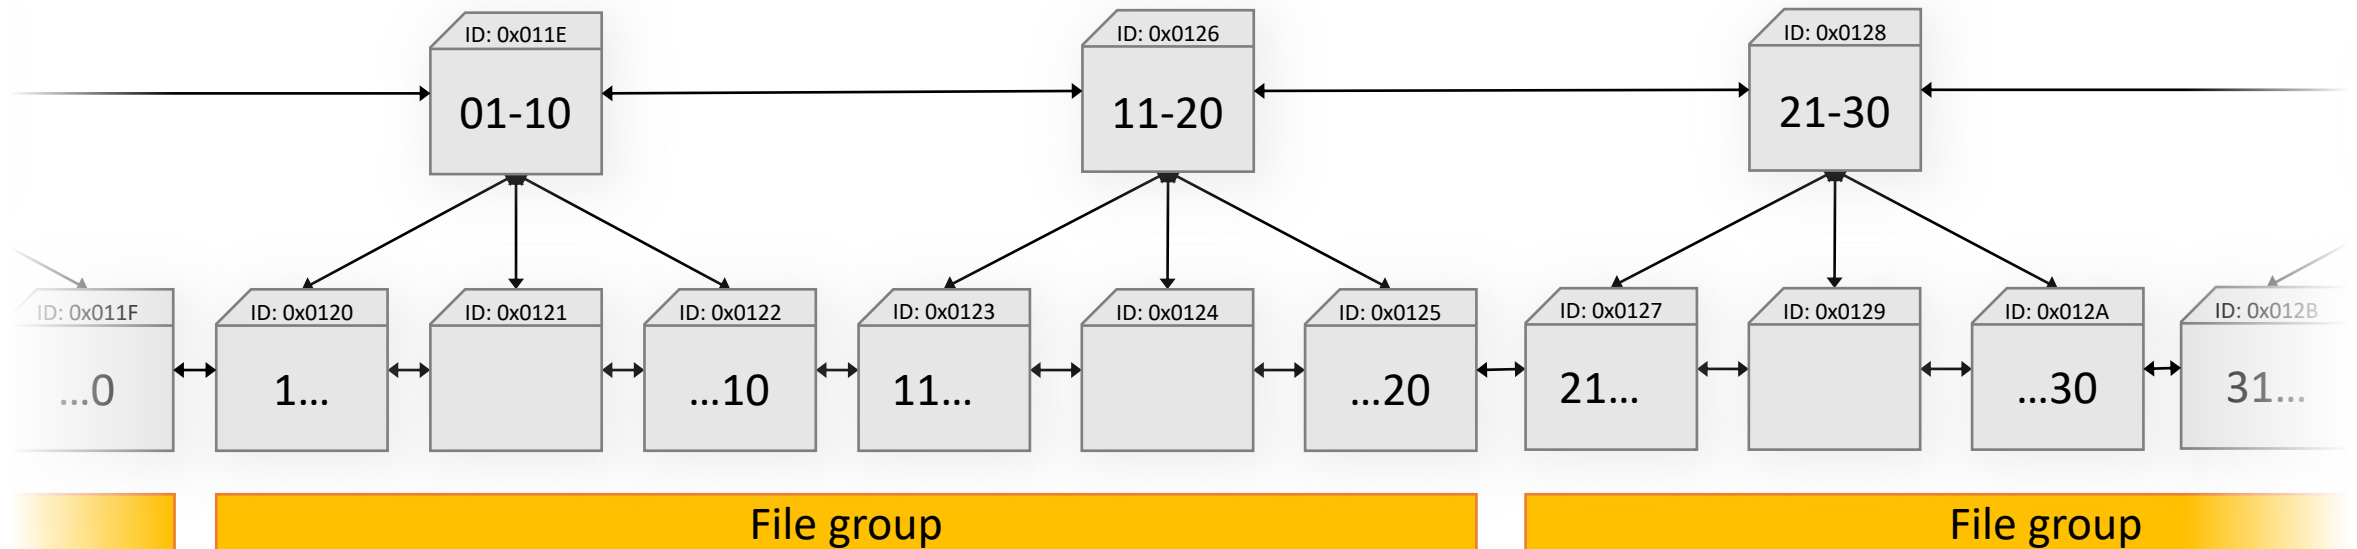




DEMO

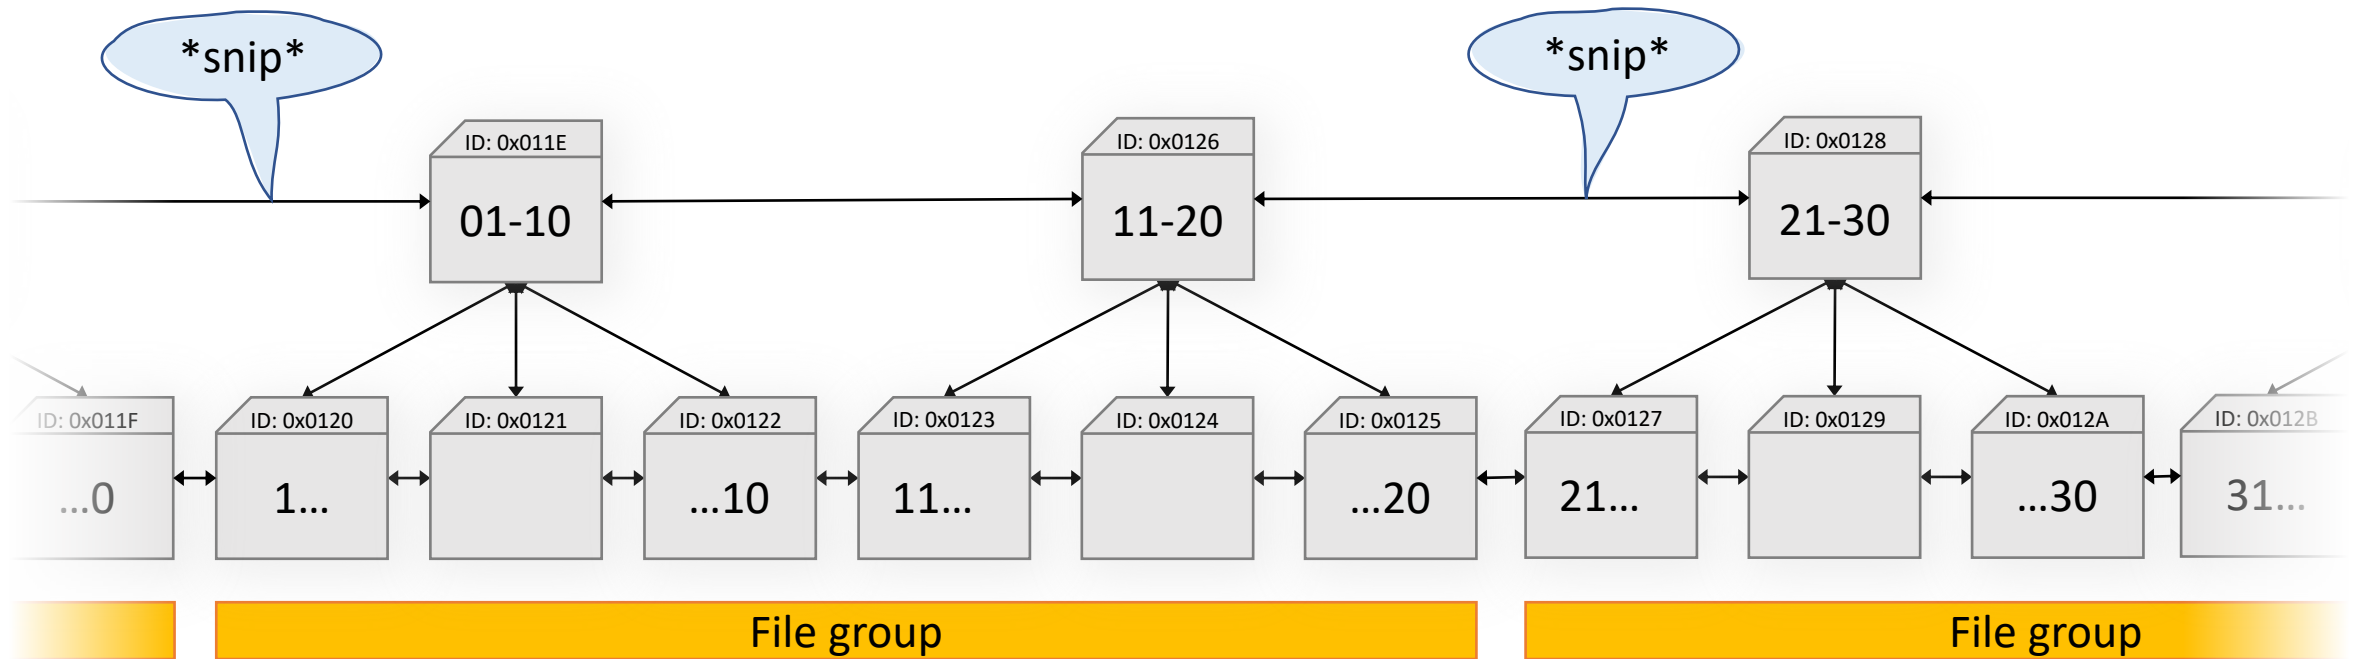
# Switching a partition

```
ALTER TABLE dbo.partitioned_table  
SWITCH PARTITION 2 TO dbo.regular_table;
```



# Switching a partition

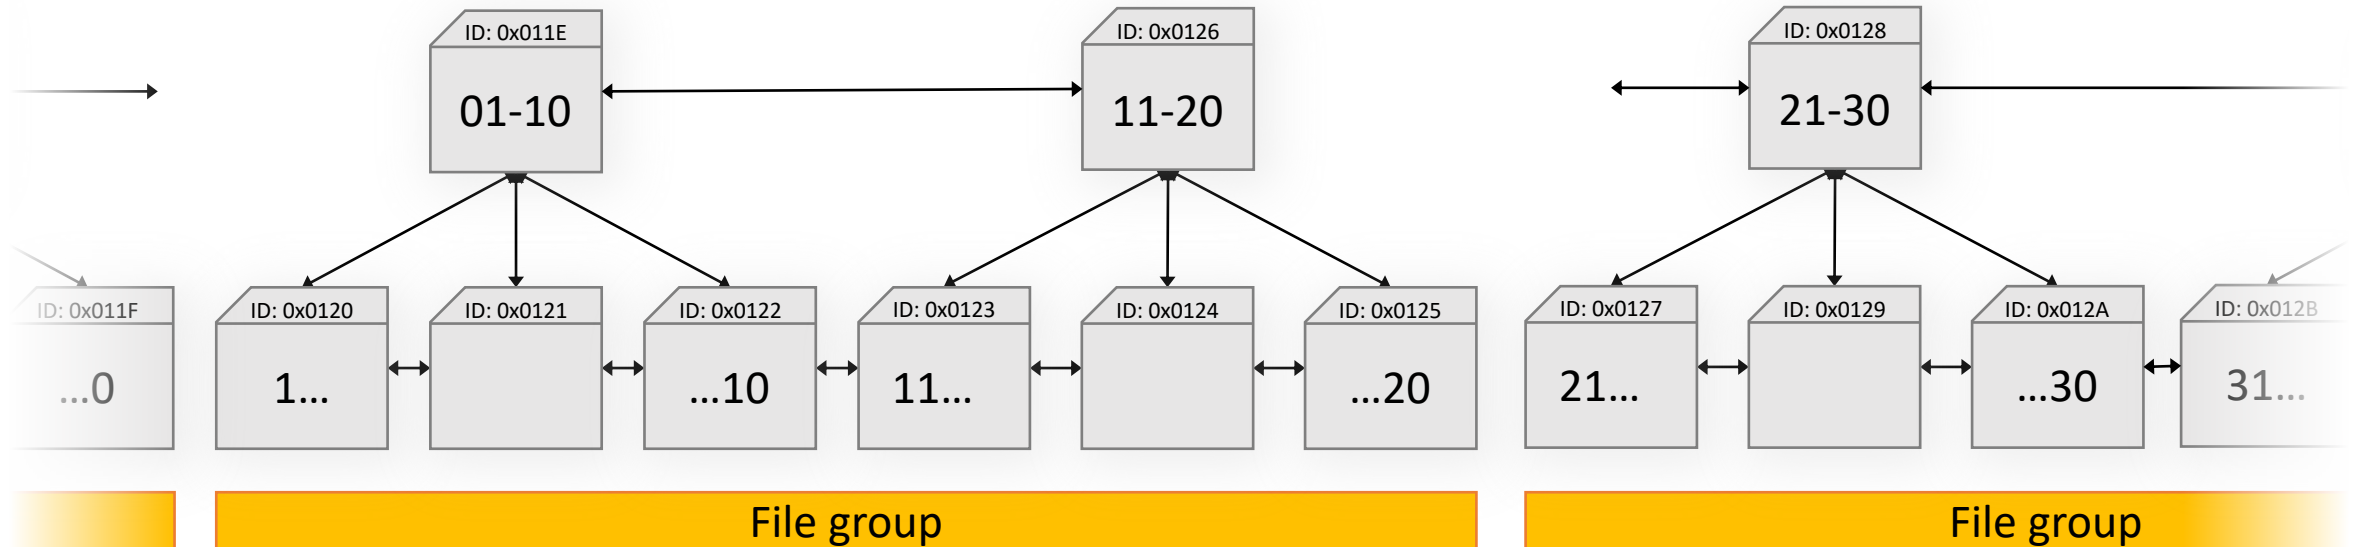
```
ALTER TABLE dbo.partitioned_table  
SWITCH PARTITION 2 TO dbo.regular_table;
```





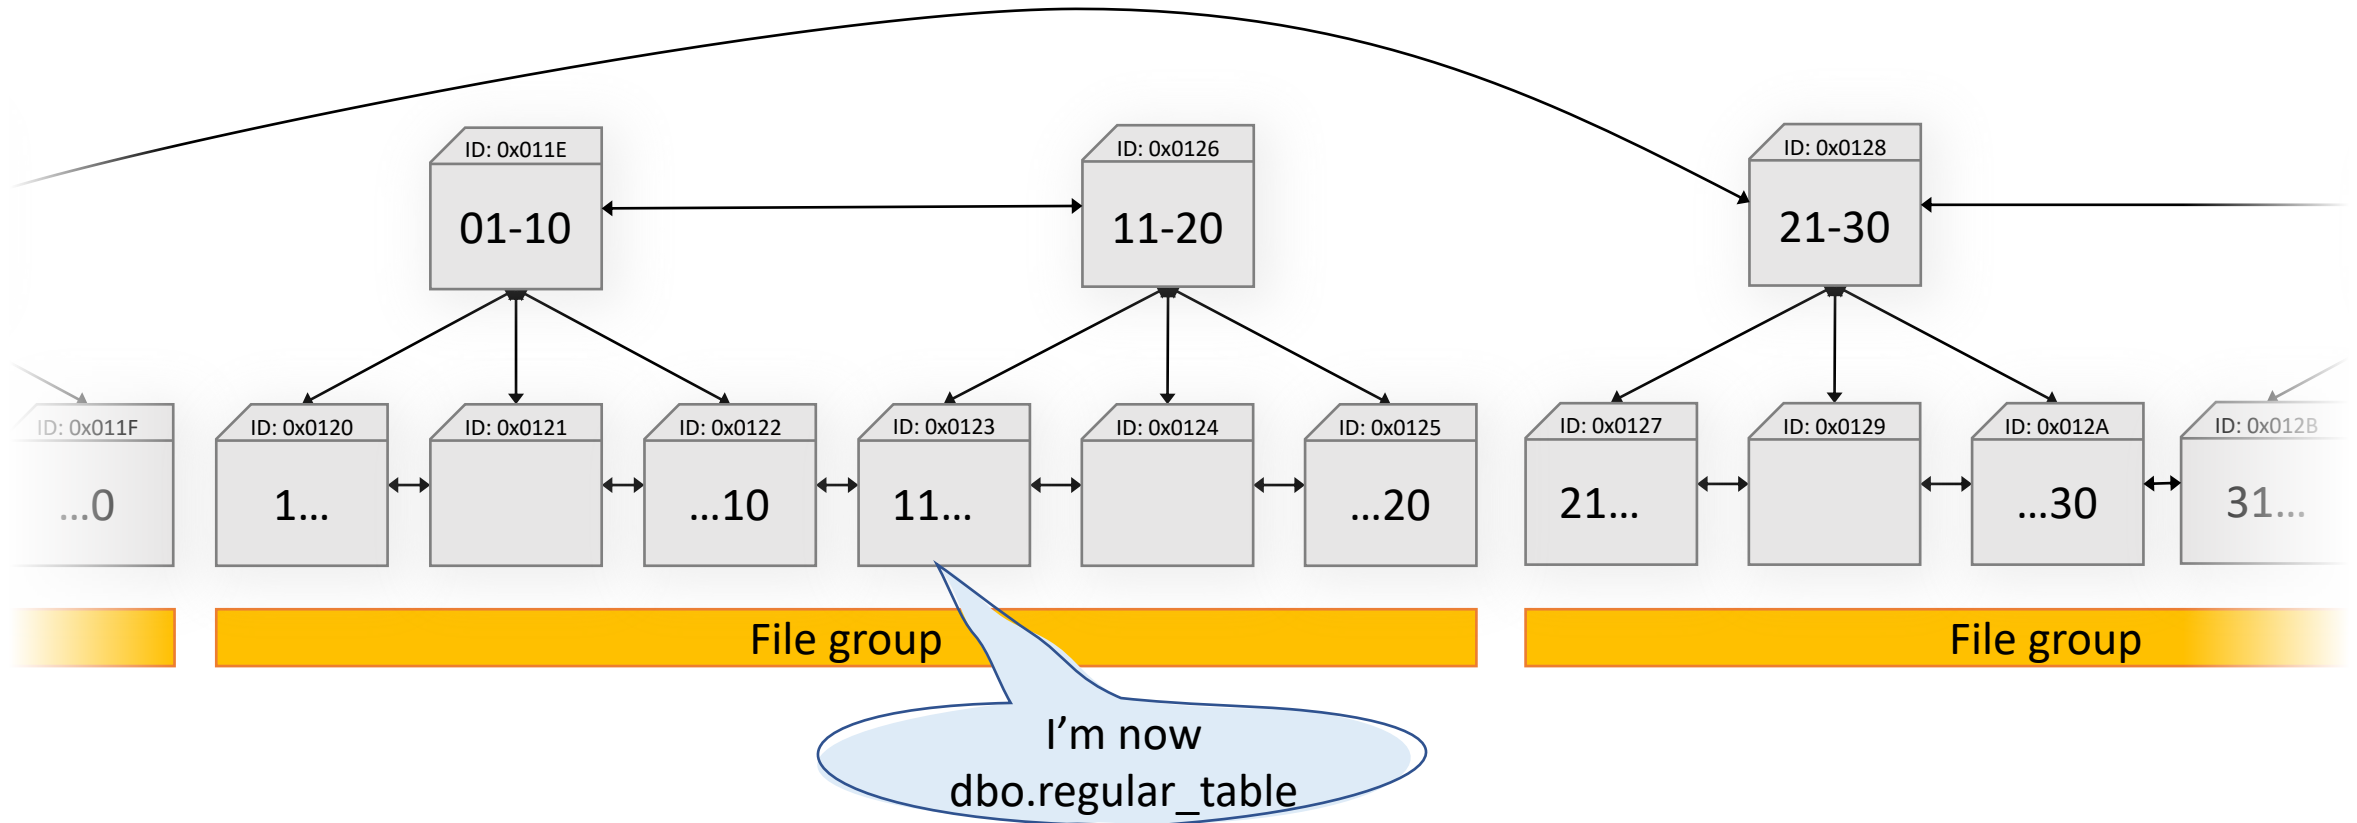
# Switching a partition

```
ALTER TABLE dbo.partitioned_table  
SWITCH PARTITION 2 TO dbo.regular_table;
```



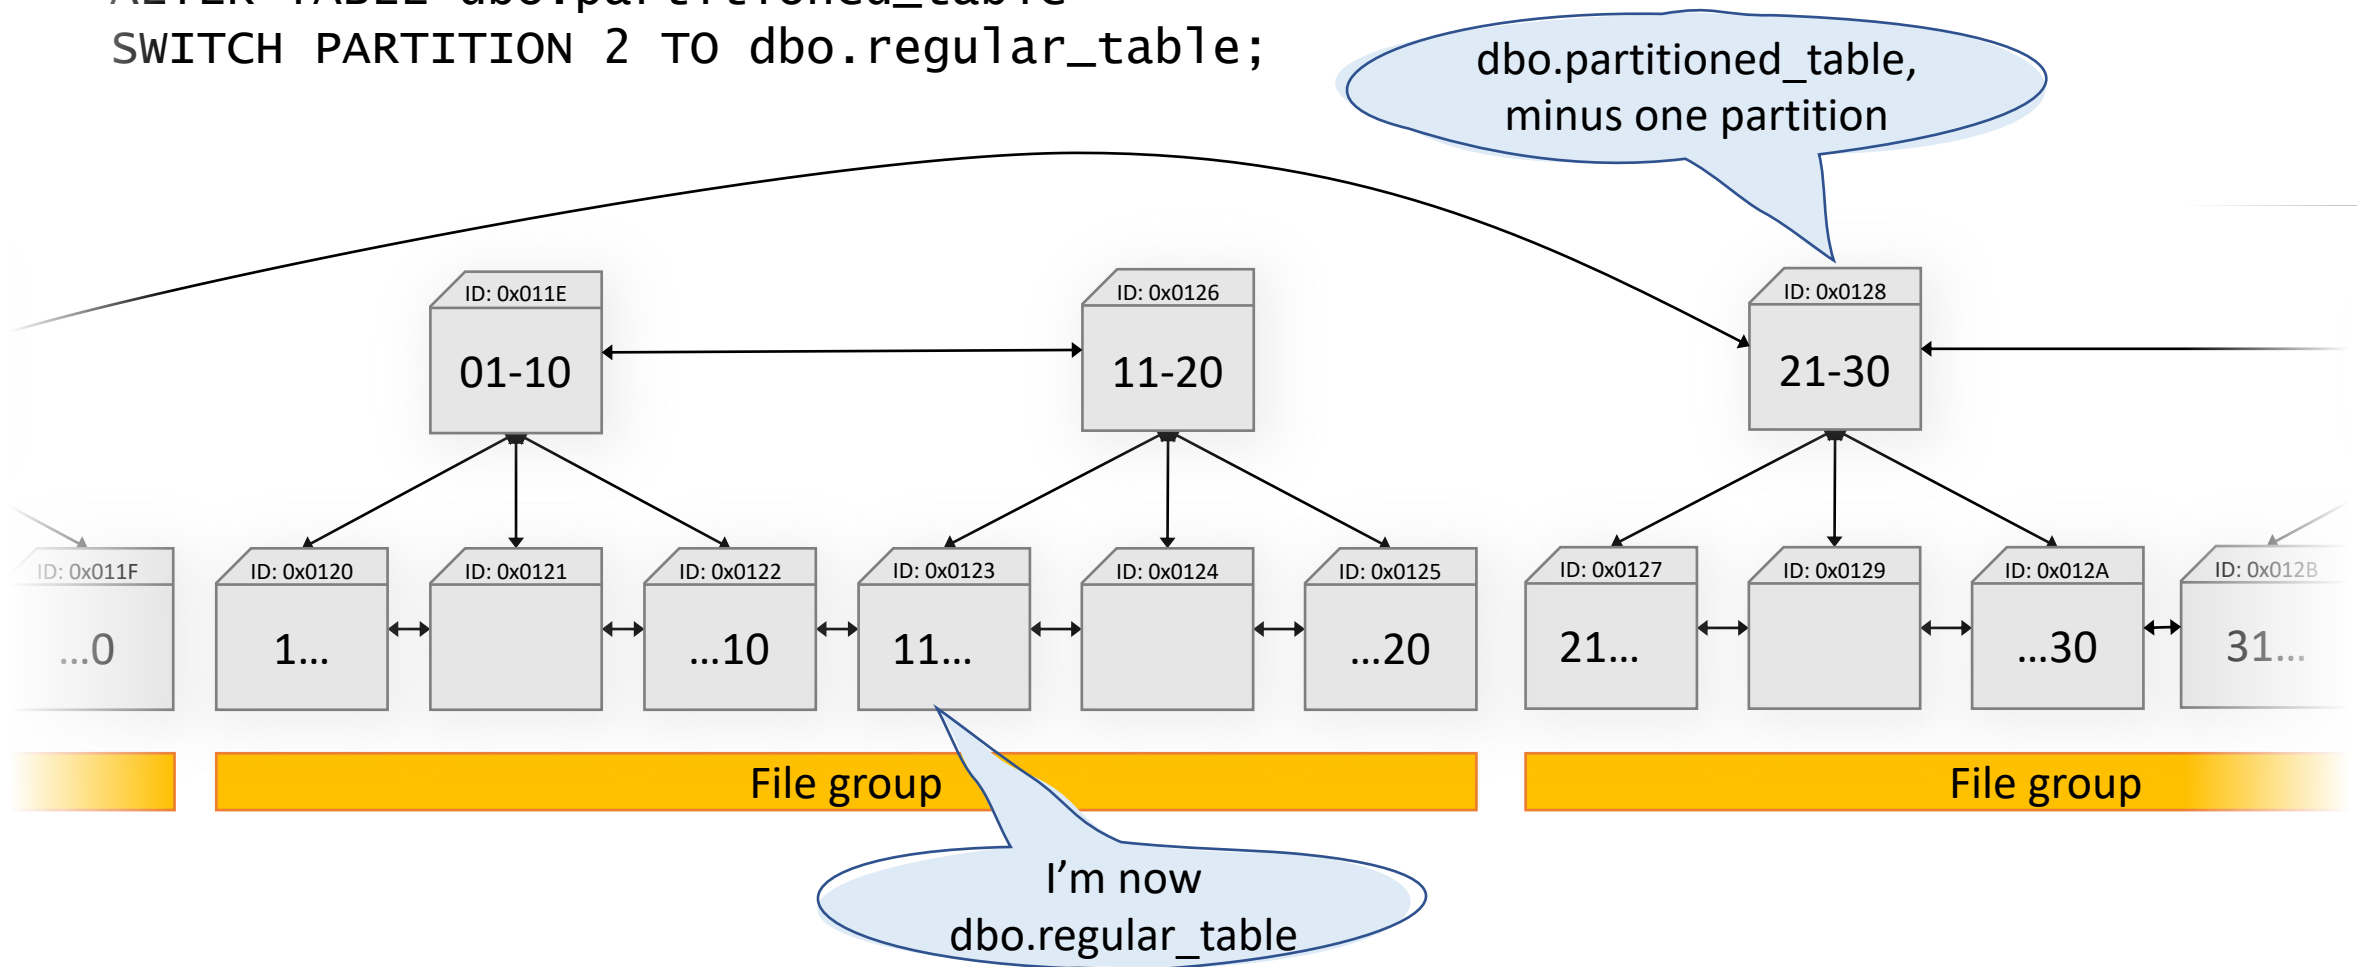
# Switching a partition

```
ALTER TABLE dbo.partitioned_table  
SWITCH PARTITION 2 TO dbo.regular_table;
```



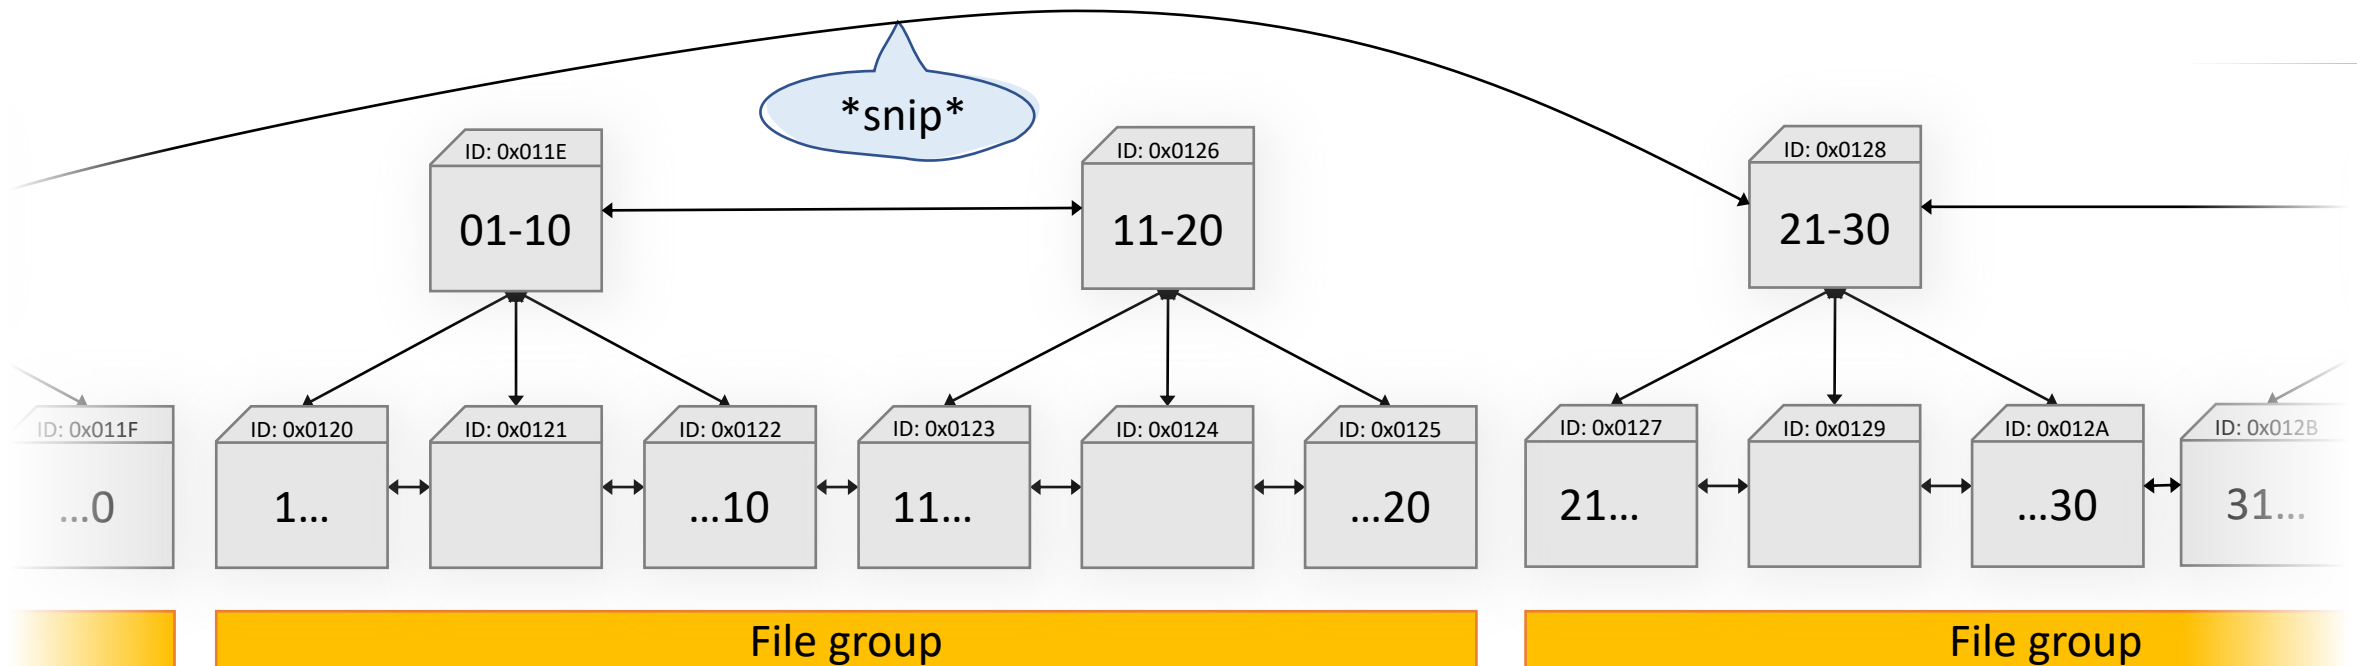
# Switching a partition

```
ALTER TABLE dbo.partitioned_table  
SWITCH PARTITION 2 TO dbo.regular_table;
```



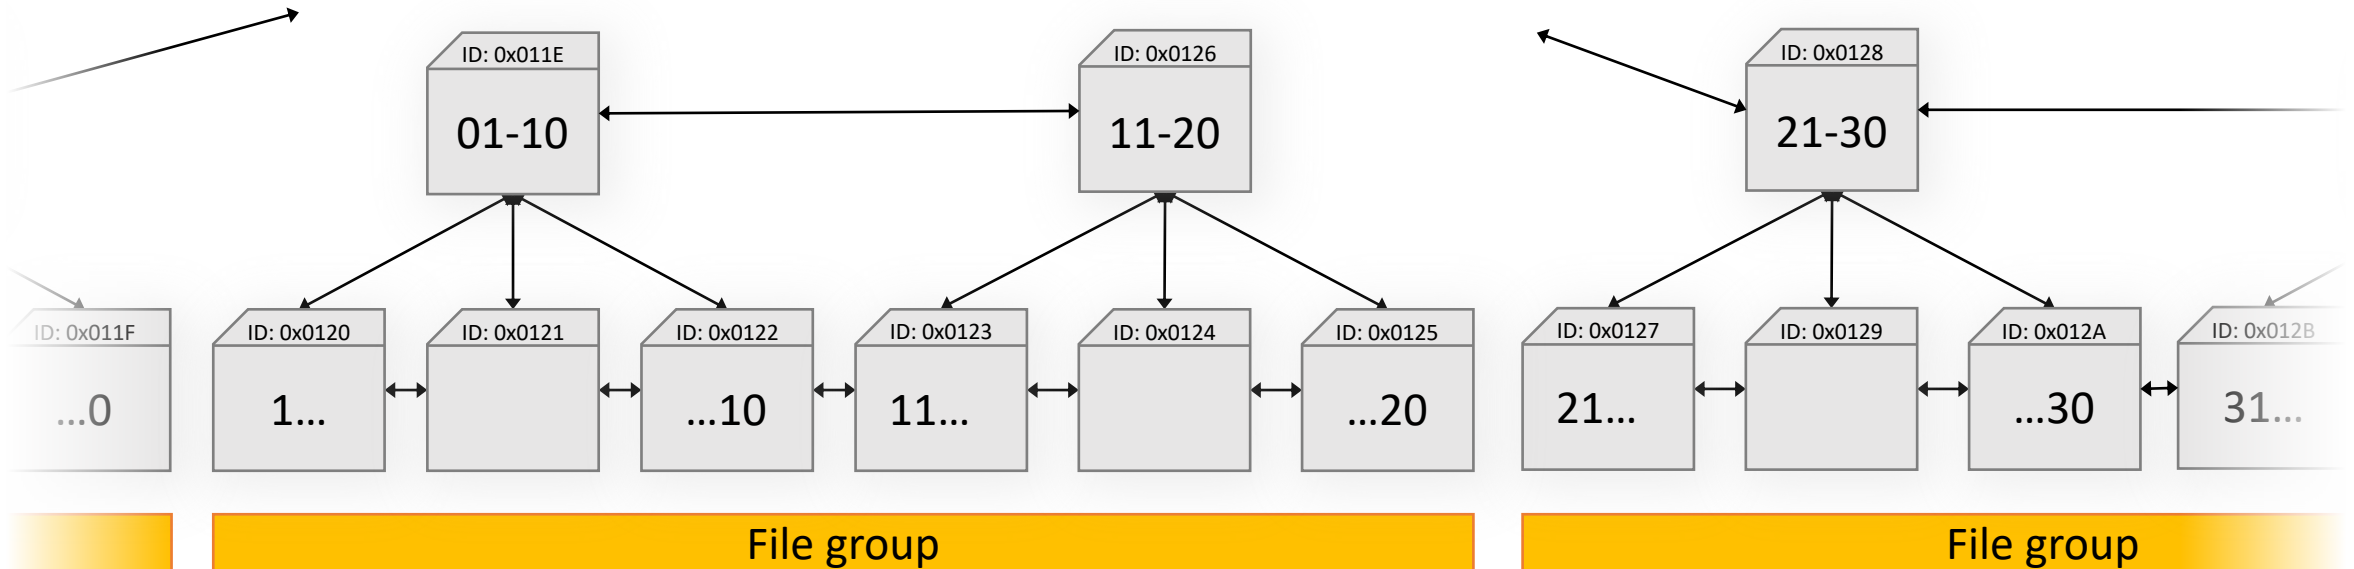
.. and back

```
ALTER TABLE dbo.regular_table  
SWITCH TO dbo.partitioned_table PARTITION 2;
```



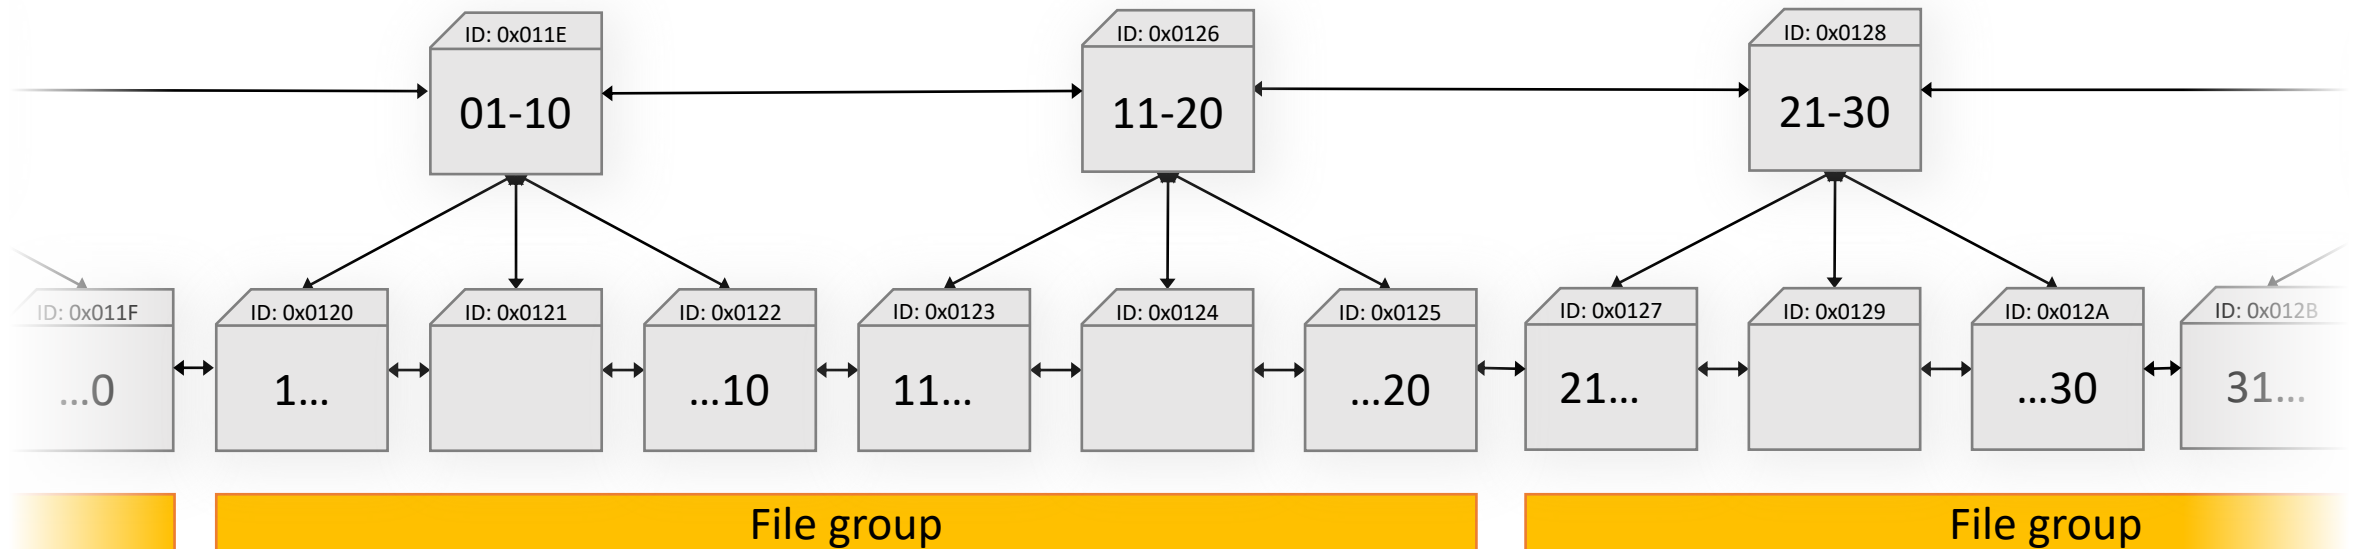
.. and back

```
ALTER TABLE dbo.regular_table  
SWITCH TO dbo.partitioned_table PARTITION 2;
```



.. and back

```
ALTER TABLE dbo.regular_table  
SWITCH TO dbo.partitioned_table PARTITION 2;
```



\* Bonus fact: you can switch two *non-partitioned* tables, too!

# These are metadata operations!

- The data doesn't move.
- Only the pointers are updated.
- Metadata operations require metadata permissions!



# DEMO



# "Right is right"

... AS RANGE LEFT  
FOR VALUES (100, 200, 300)

	$n \leq$	100
100	$< n \leq$	200
200	$< n \leq$	300
300	$< n$	

---

... AS RANGE RIGHT  
FOR VALUES (100, 200, 300)

	$n <$	100
100	$\leq n <$	200
200	$\leq n <$	300
300	$\leq n$	

---

# "Right is right"

```
CREATE PARTITION FUNCTION AnnualFunction(date)
AS RANGE RIGHT
FOR VALUES ( '2020-01-01' , '2021-01-01' , '2022-01-01' );
```

```
CREATE PARTITION SCHEME Annual
AS PARTITION AnnualFunction
TO ( [FILEGROUP_A] , [FILEGROUP_B] , [FILEGROUP_C] , [FILEGROUP_D] );
```

< 2020-01-01

≥ 2020-01-01  
< 2021-01-01

# "Right is right"

```
CREATE PARTITION FUNCTION AnnualFunction(date)
```

```
AS RANGE LEFT
```

```
FOR VALUES ( '2020-01-01' , '2021-01-01' , '2022-01-01' );
```

```
CREATE PARTITION SCHEME Annual
```

```
AS PARTITION AnnualFunction
```

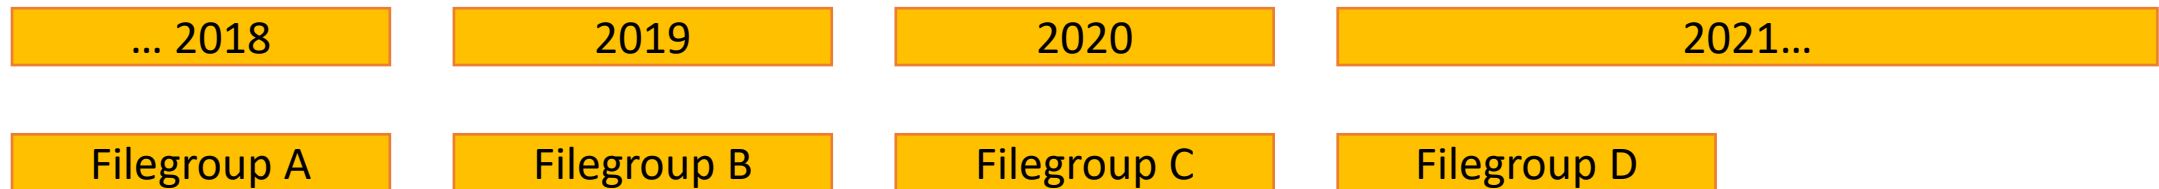
```
TO ( [FILEGROUP_A] , [FILEGROUP_B] , [FILEGROUP_C] , [..._D] );
```

$\leq 2020-01-01$

$> 2020-01-01$   
 $\leq 2021-01-01$

# Split/merge

- You can *split* an existing partition into two, by adding a new boundary.



# Split/merge

- You can *split* an existing partition into two, by adding a new boundary.



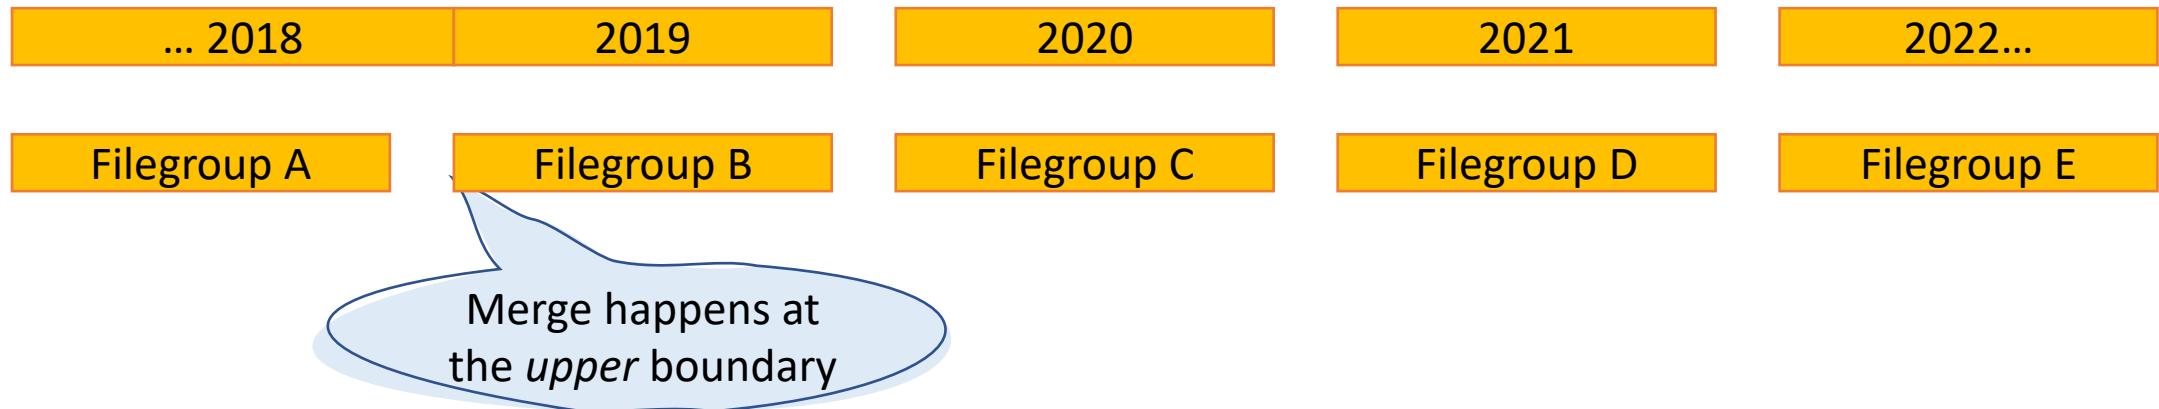
# Split/merge

- You can *split* an existing partition into two, by adding a new boundary.



# Split/merge

- You can *split* an existing partition into two, by adding a new boundary.
- Two partitions can also be *merged* into one, by removing a boundary.

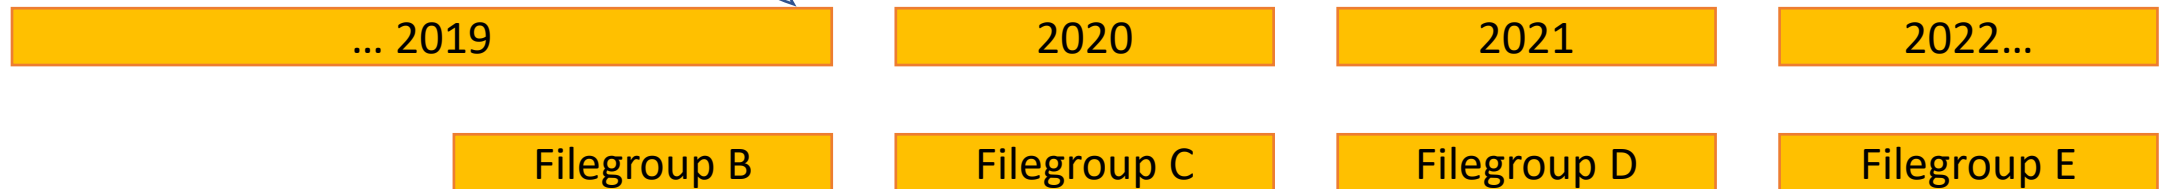




# Split/merge

- You can *split* an existing partition into two, by adding a new boundary.
- Two partitions can also be *merged* into one, by removing a boundary.

.. so the merged partition stays on filegroup B.



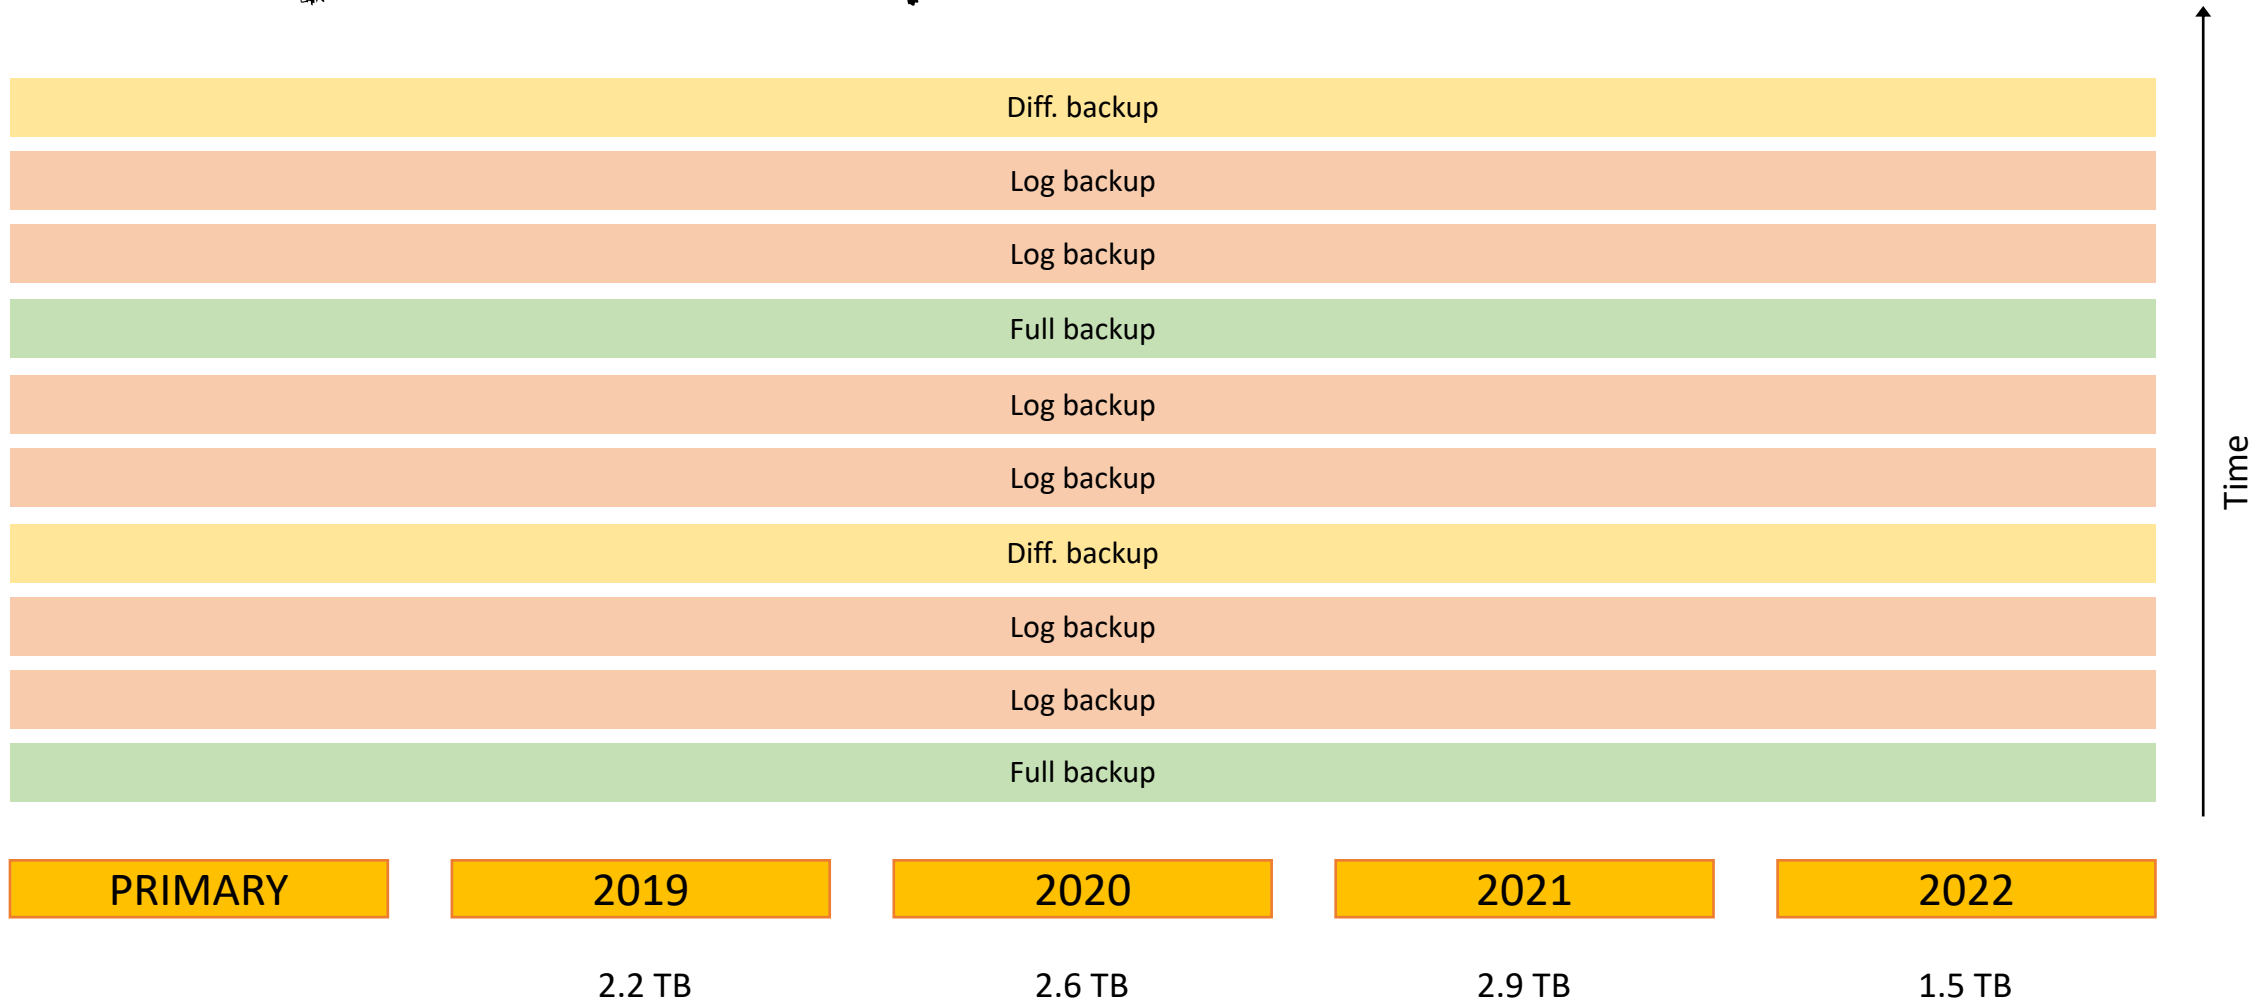


DEMO

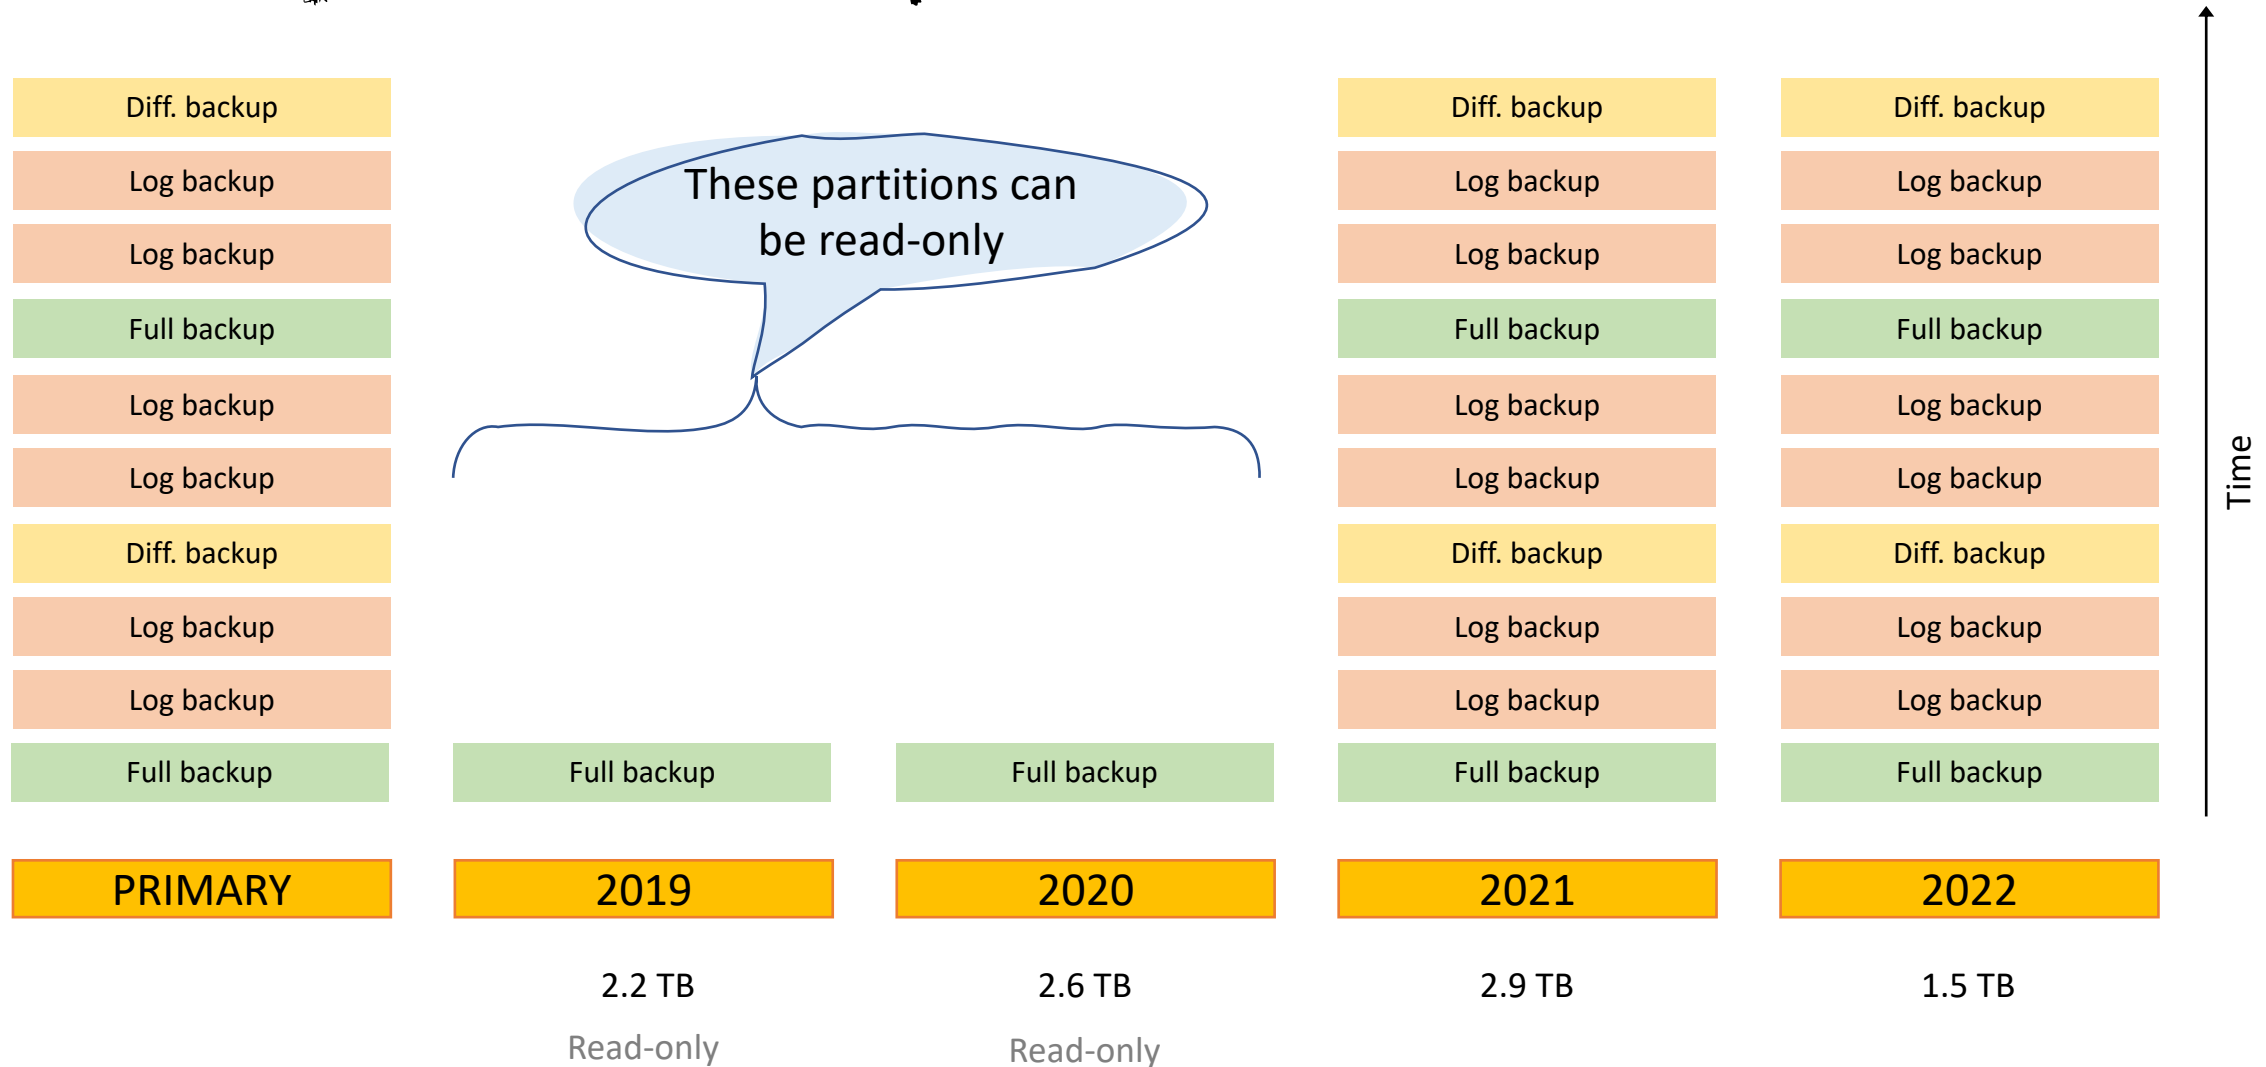
# Aligned / non-aligned partitions

- If the table and its indexes all share the same partition scheme, the indexes are “aligned”
- If one or more indexes are not partitioned, or differently partitioned, they are “non-aligned”, and you won’t be able to perform partition-level operations like SPLIT, MERGE, TRUNCATE, SWITCH.

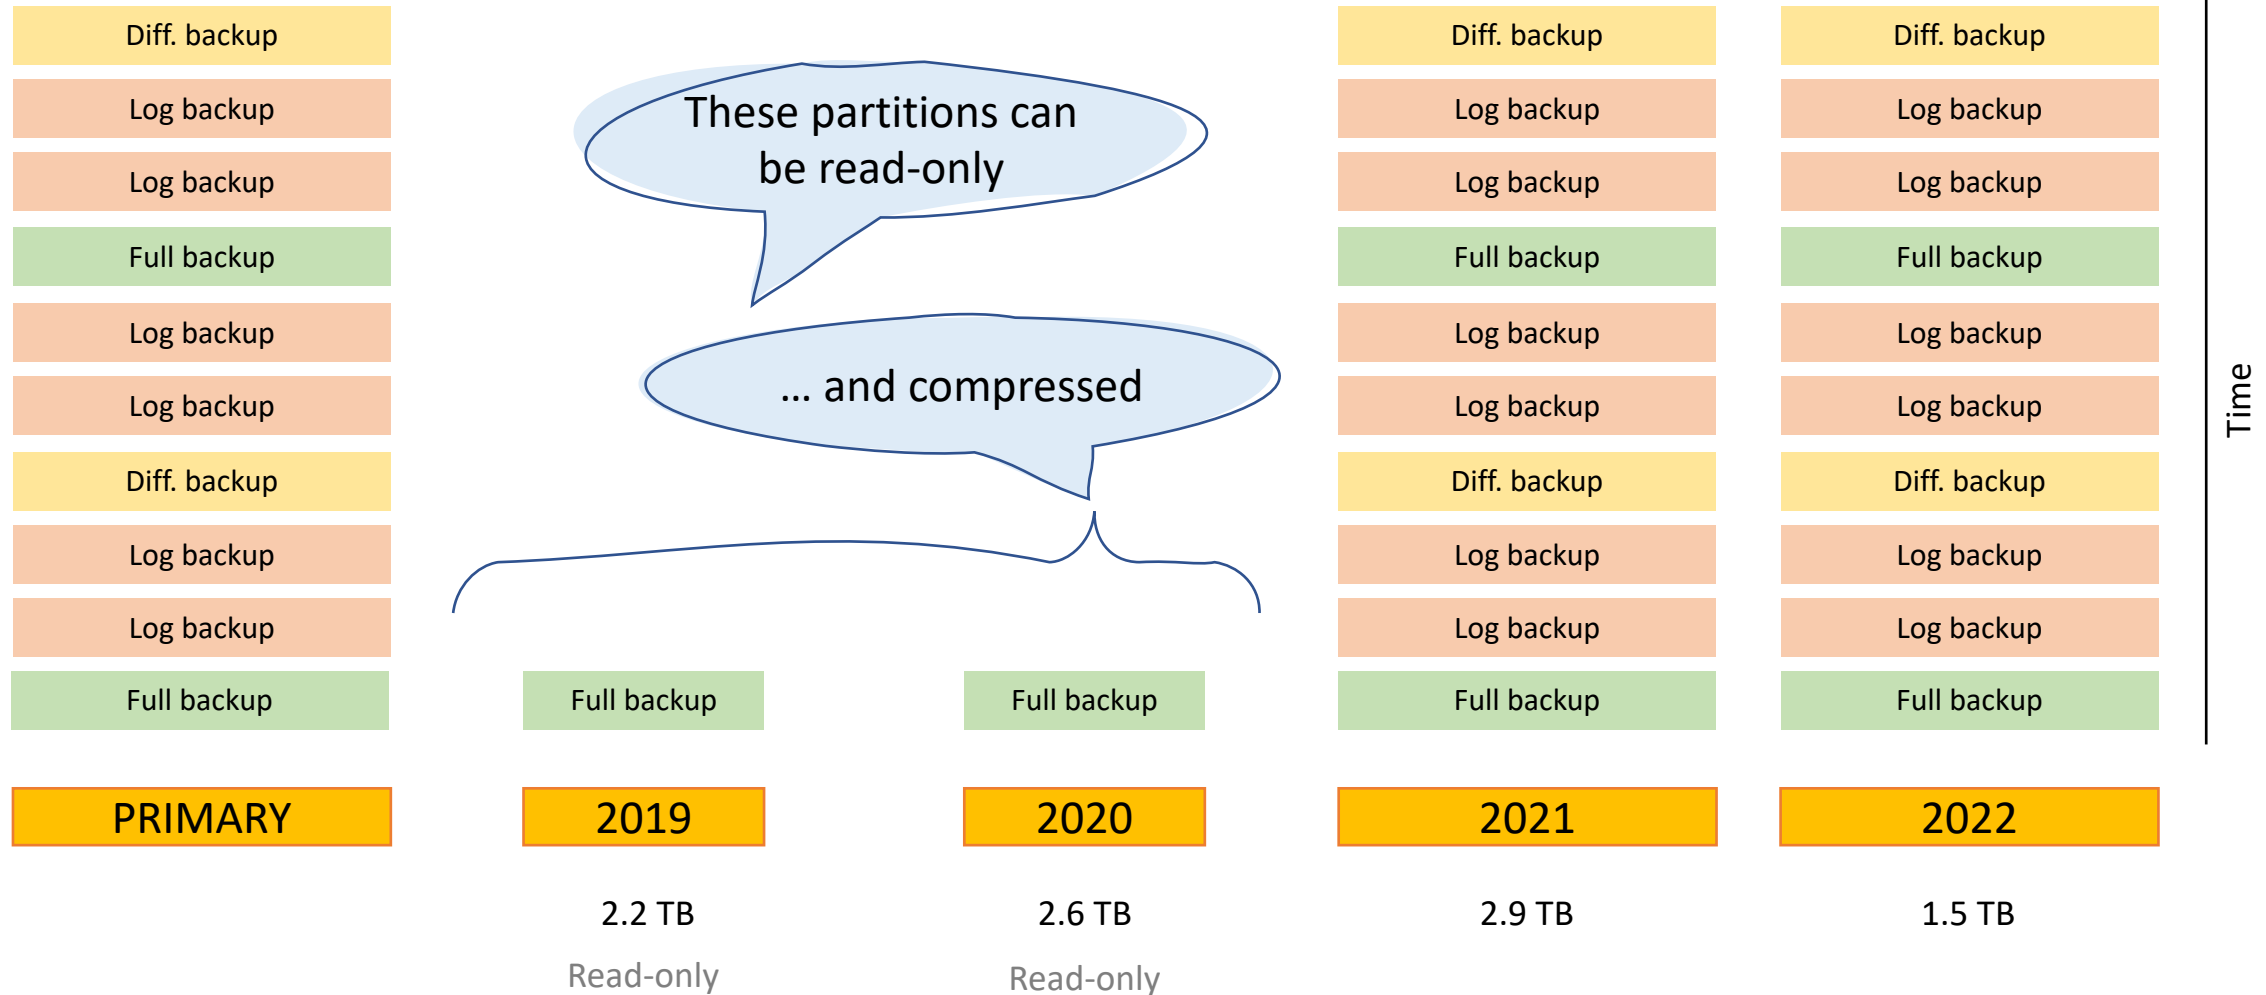
# Make your backups smaller



# Make your backups smaller



# Make your backups smaller



# Piecemeal restore (full)

We have:

- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups



# Piecemeal restore (full)

We have:

- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups

Full backup, partial,  
recovery

PRIMARY

Some file

2019

Read-only

2020

Read-only

2021

2022

2023



## Piecemeal restore (full)

We have:

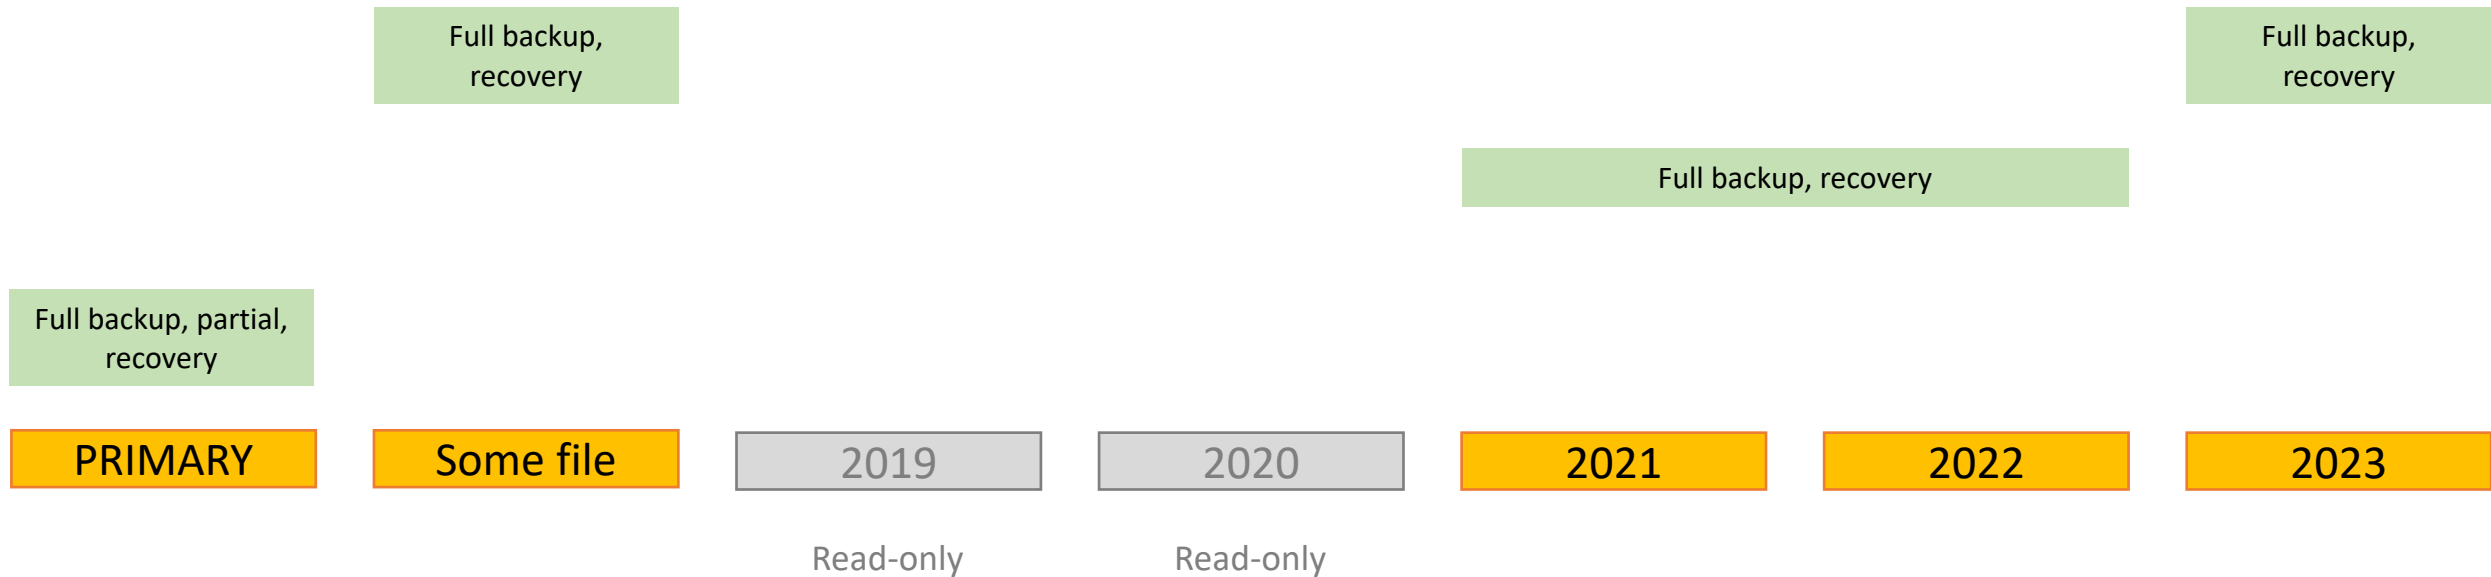
- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups



# Piecemeal restore (full)

We have:

- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups

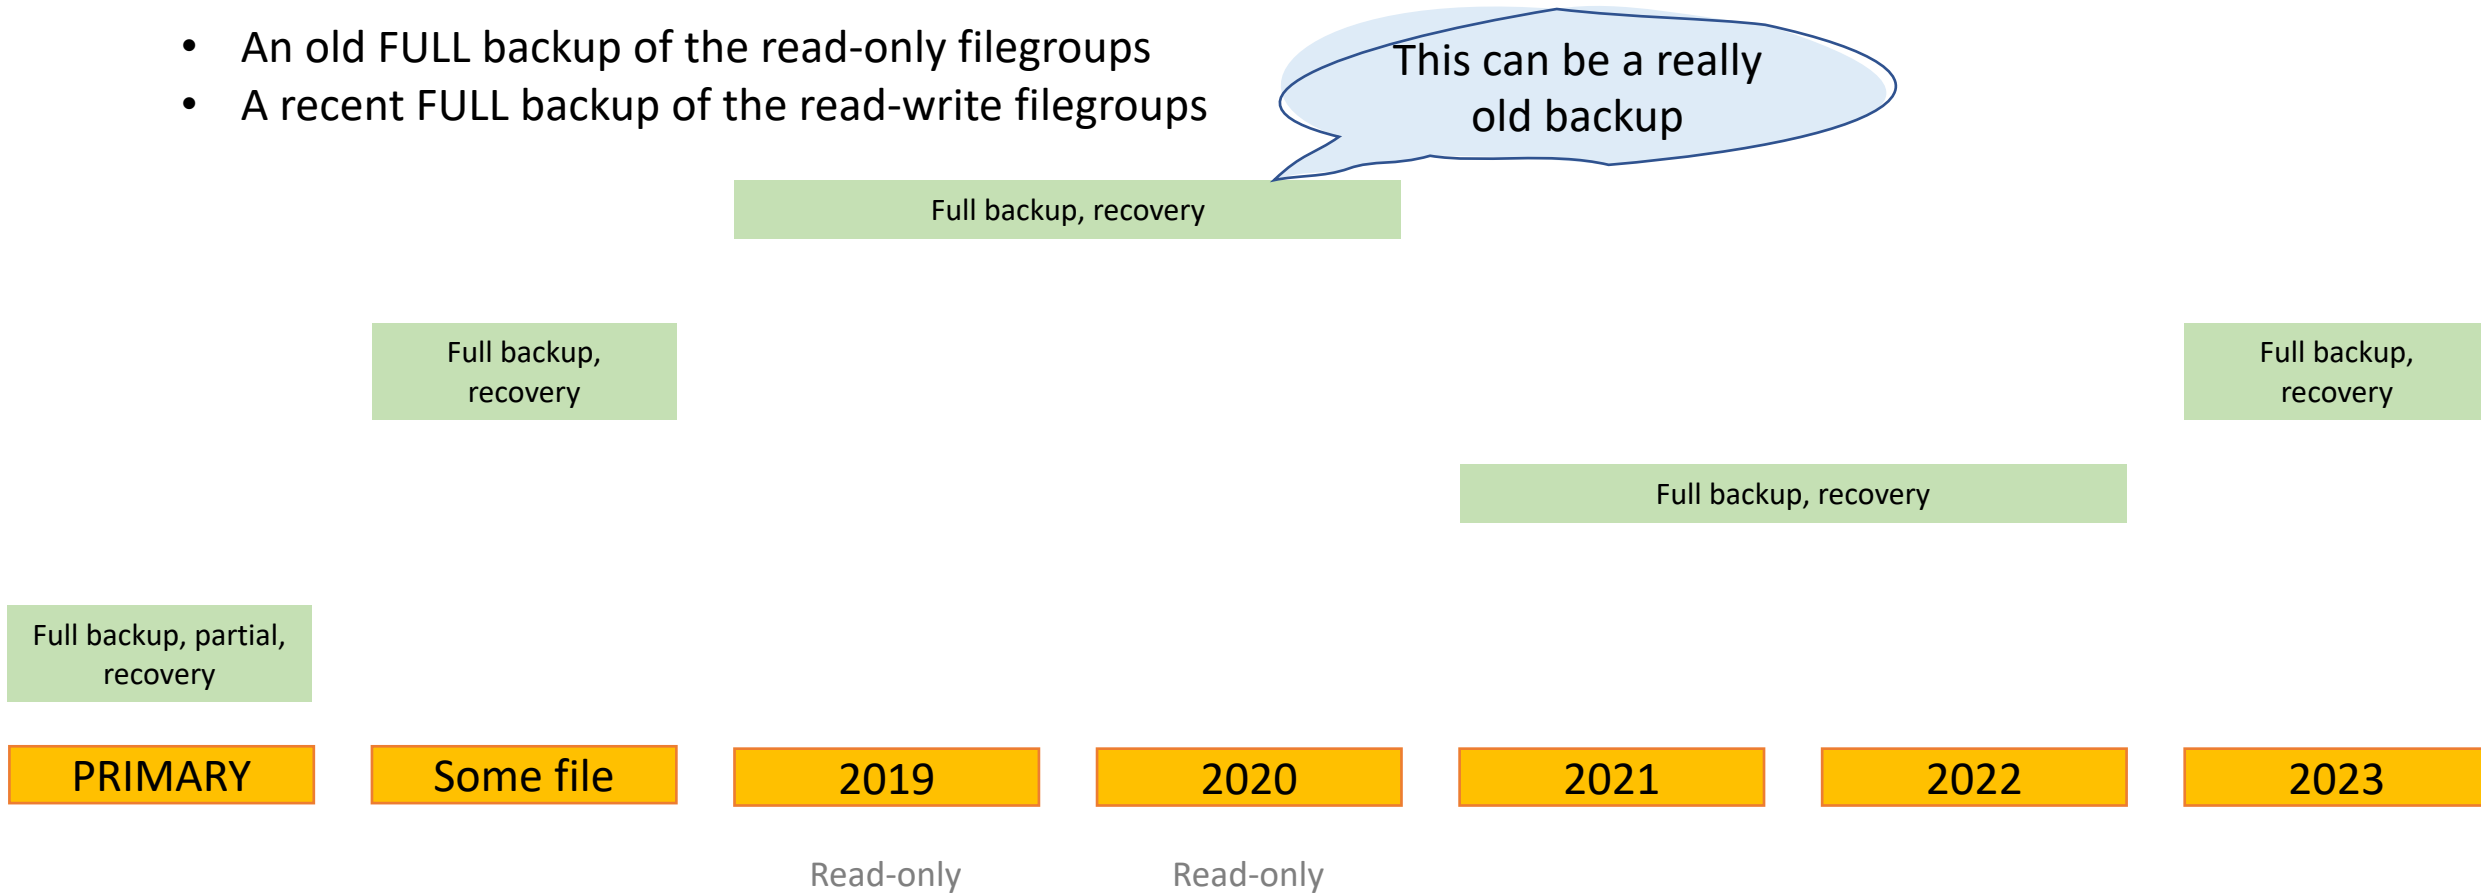


# Piecemeal restore (full)

We have:

- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups

This can be a really old backup







SSC

DEMO



# Piecemeal restore (full + log)

We have:

- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups
- Transaction log backup(s) for the read-write filegroups



# Piecemeal restore (full + log)

We have:

- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups
- Transaction log backup(s) for the read-write filegroups

Full, partial, norec.

PRIMARY

Some file

2019

Read-only

2020

Read-only

2021

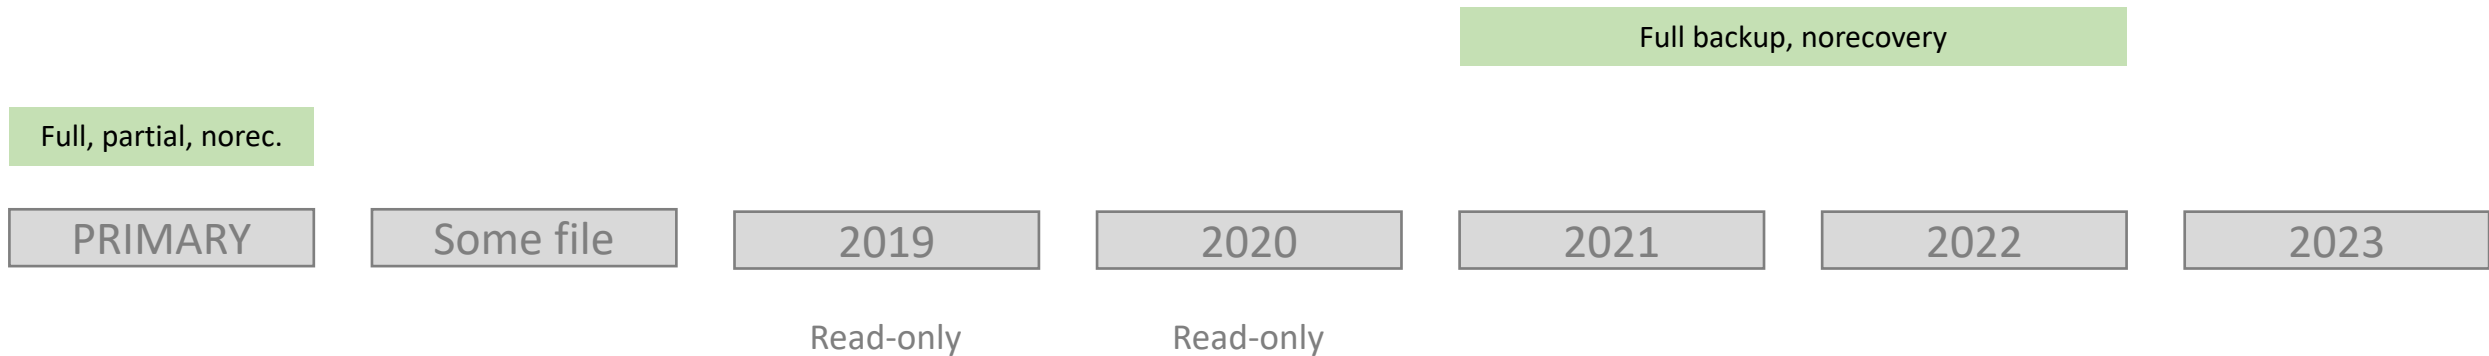
2022

2023

# Piecemeal restore (full + log)

We have:

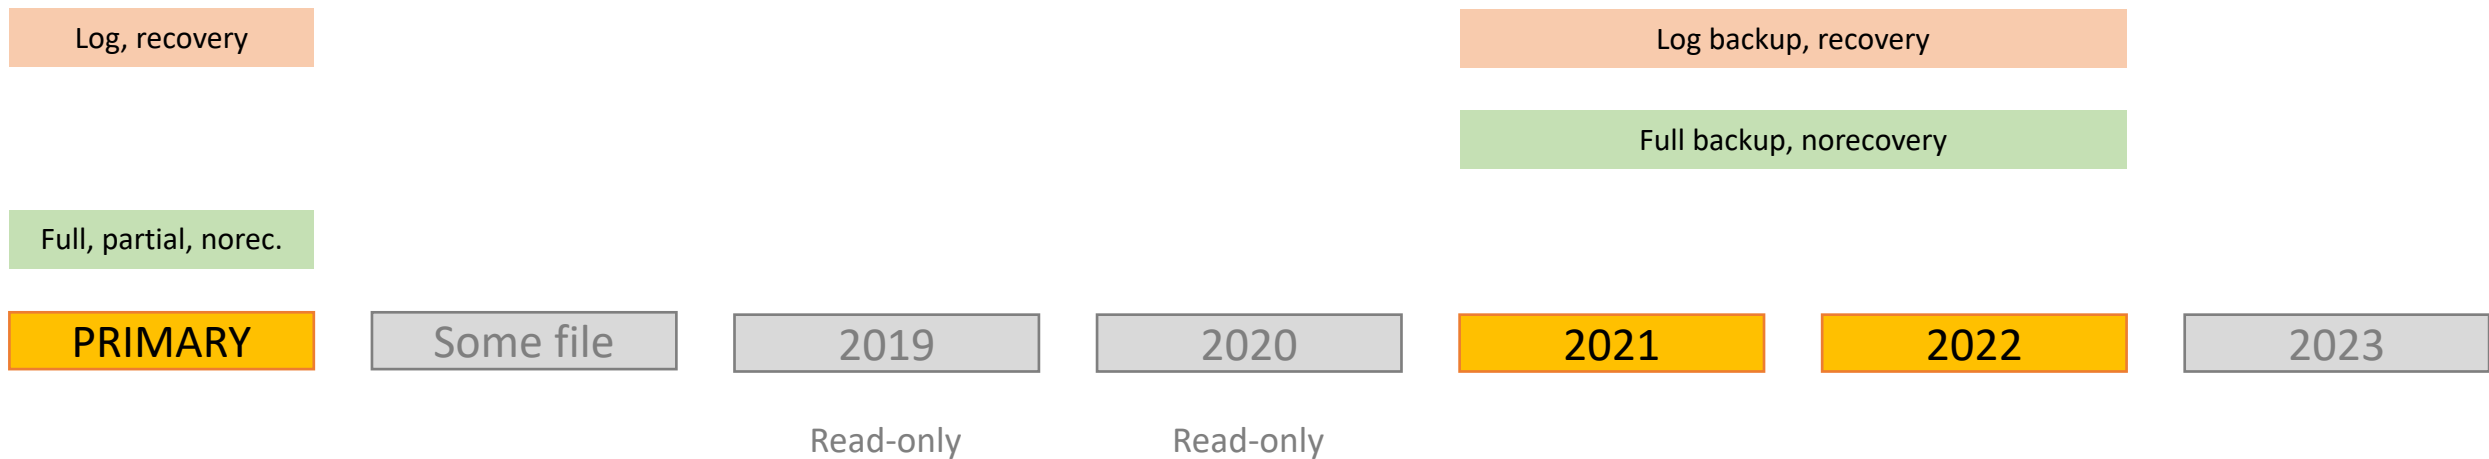
- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups
- Transaction log backup(s) for the read-write filegroups



# Piecemeal restore (full + log)

We have:

- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups
- Transaction log backup(s) for the read-write filegroups

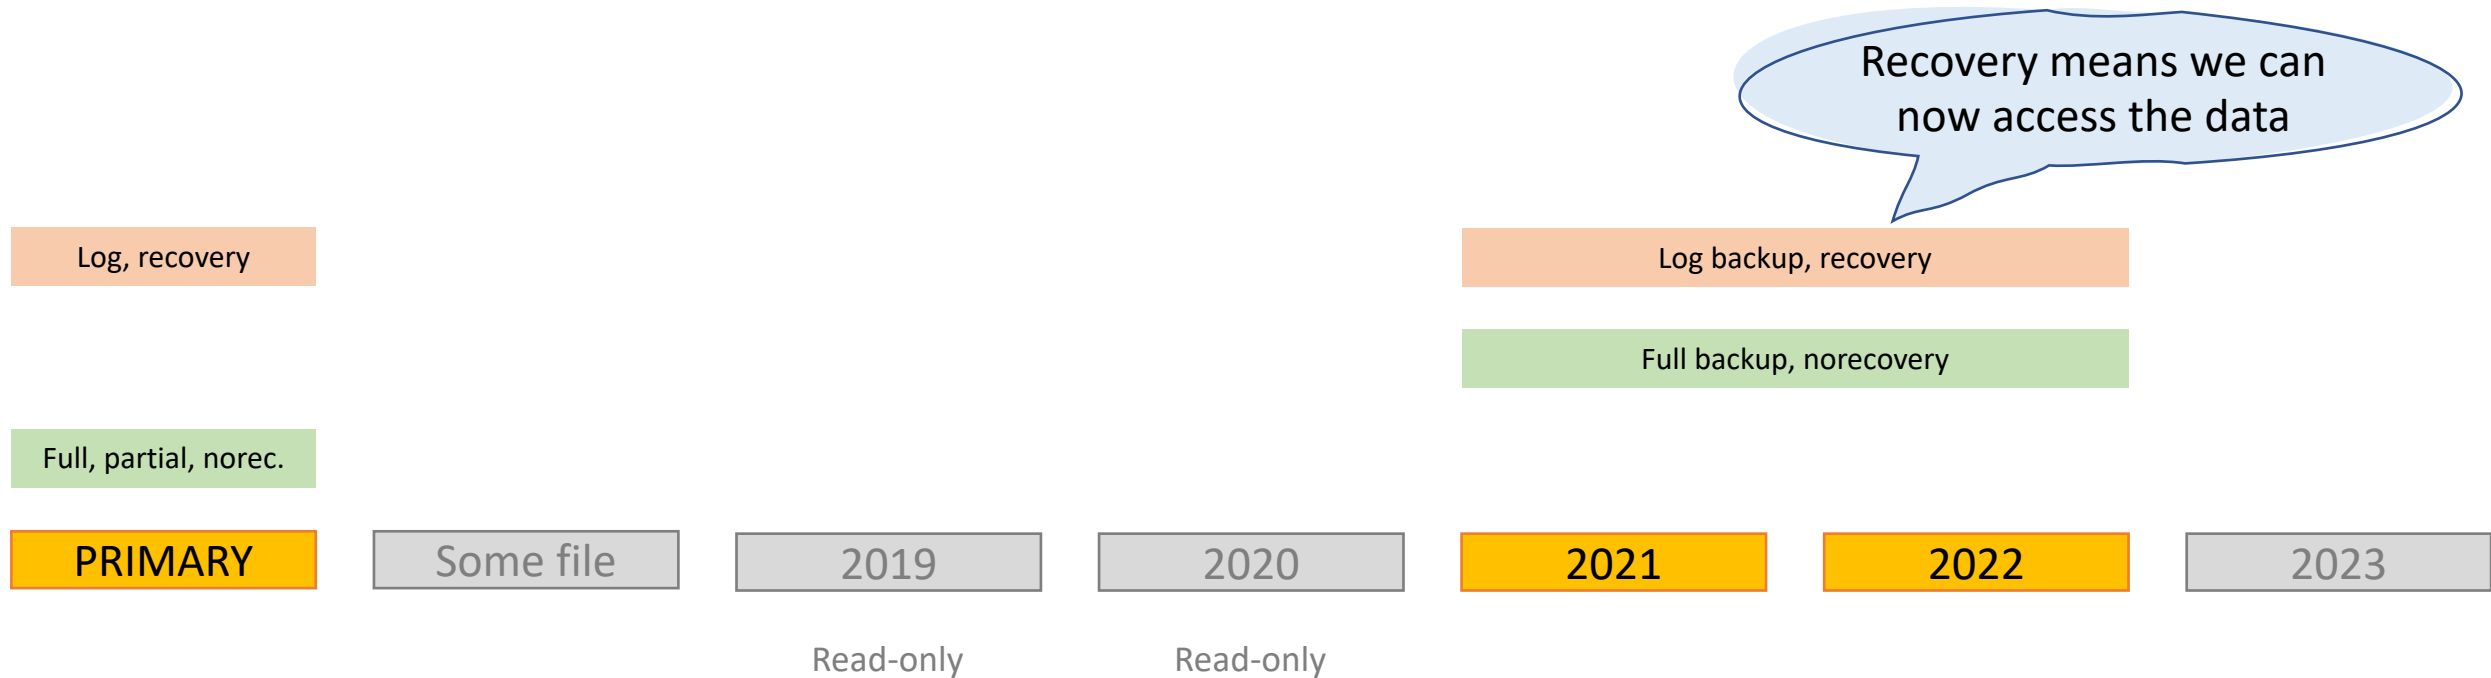




# Piecemeal restore (full + log)

We have:

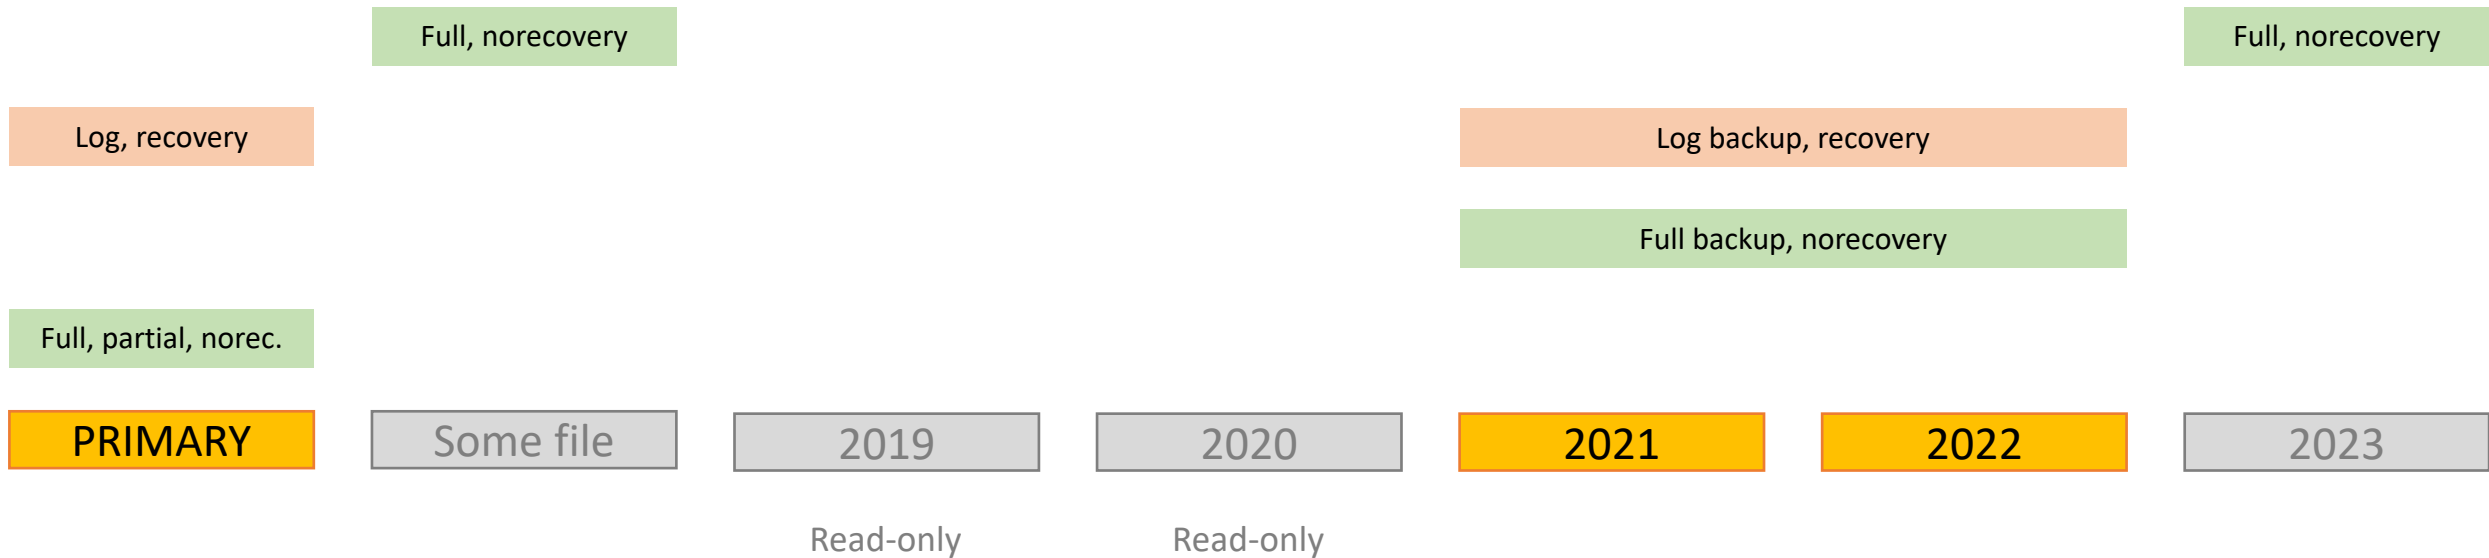
- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups
- Transaction log backup(s) for the read-write filegroups



# Piecemeal restore (full + log)

We have:

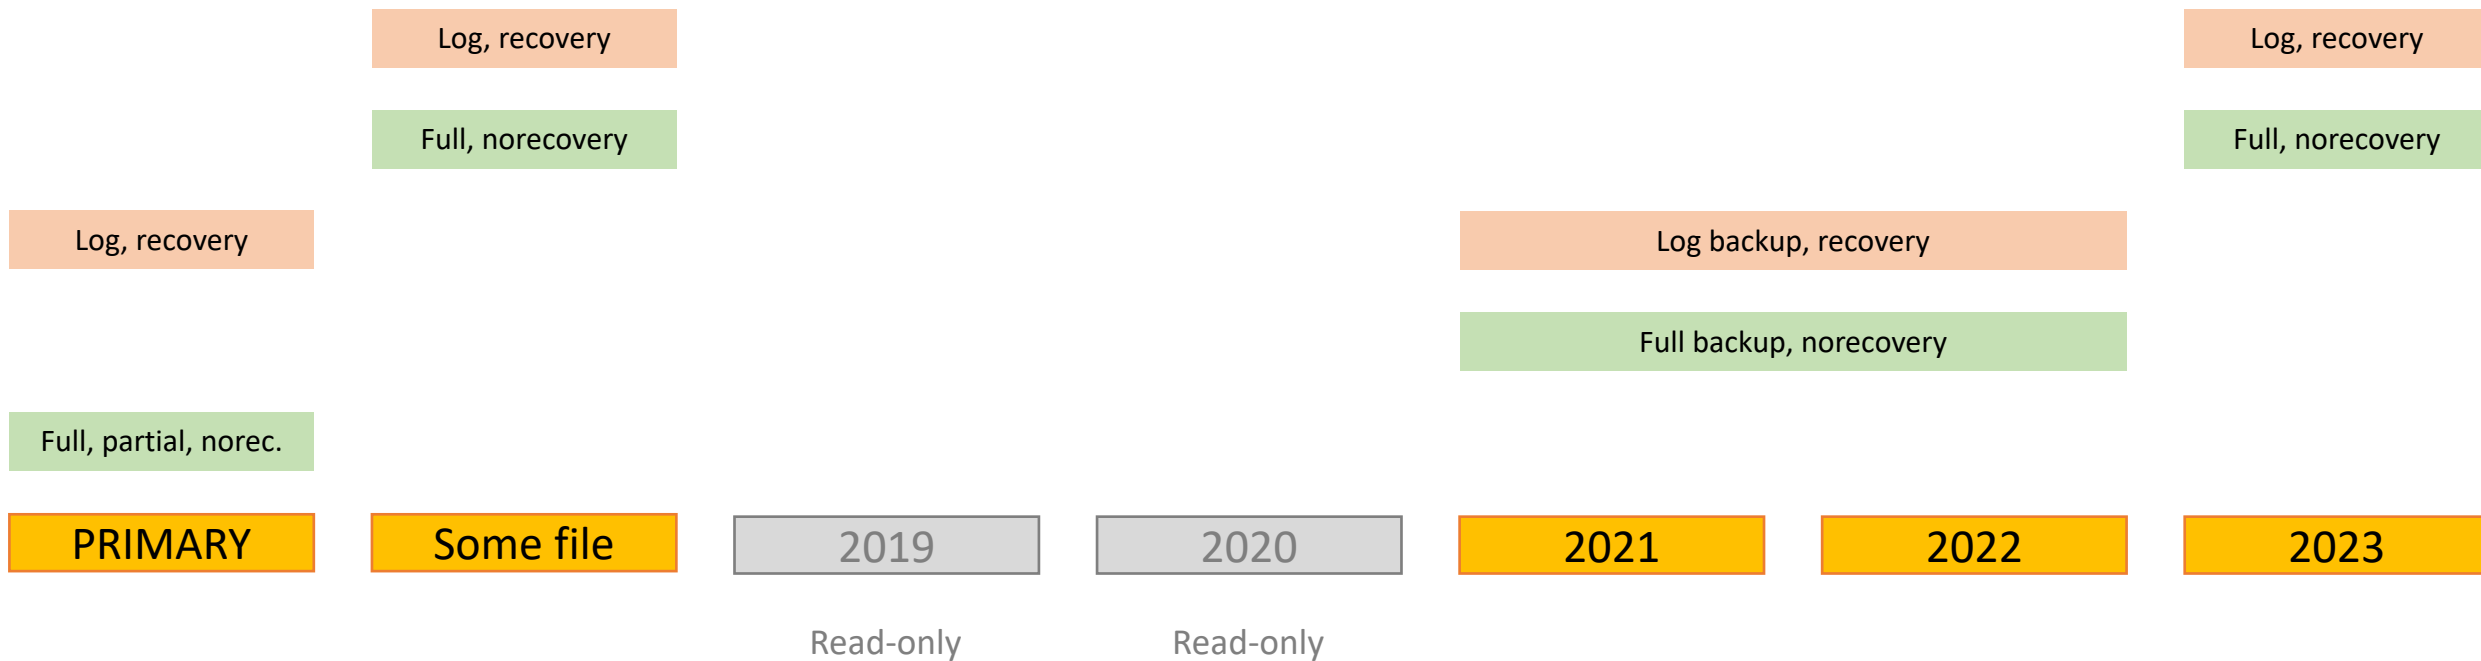
- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups
- Transaction log backup(s) for the read-write filegroups



# Piecemeal restore (full + log)

We have:

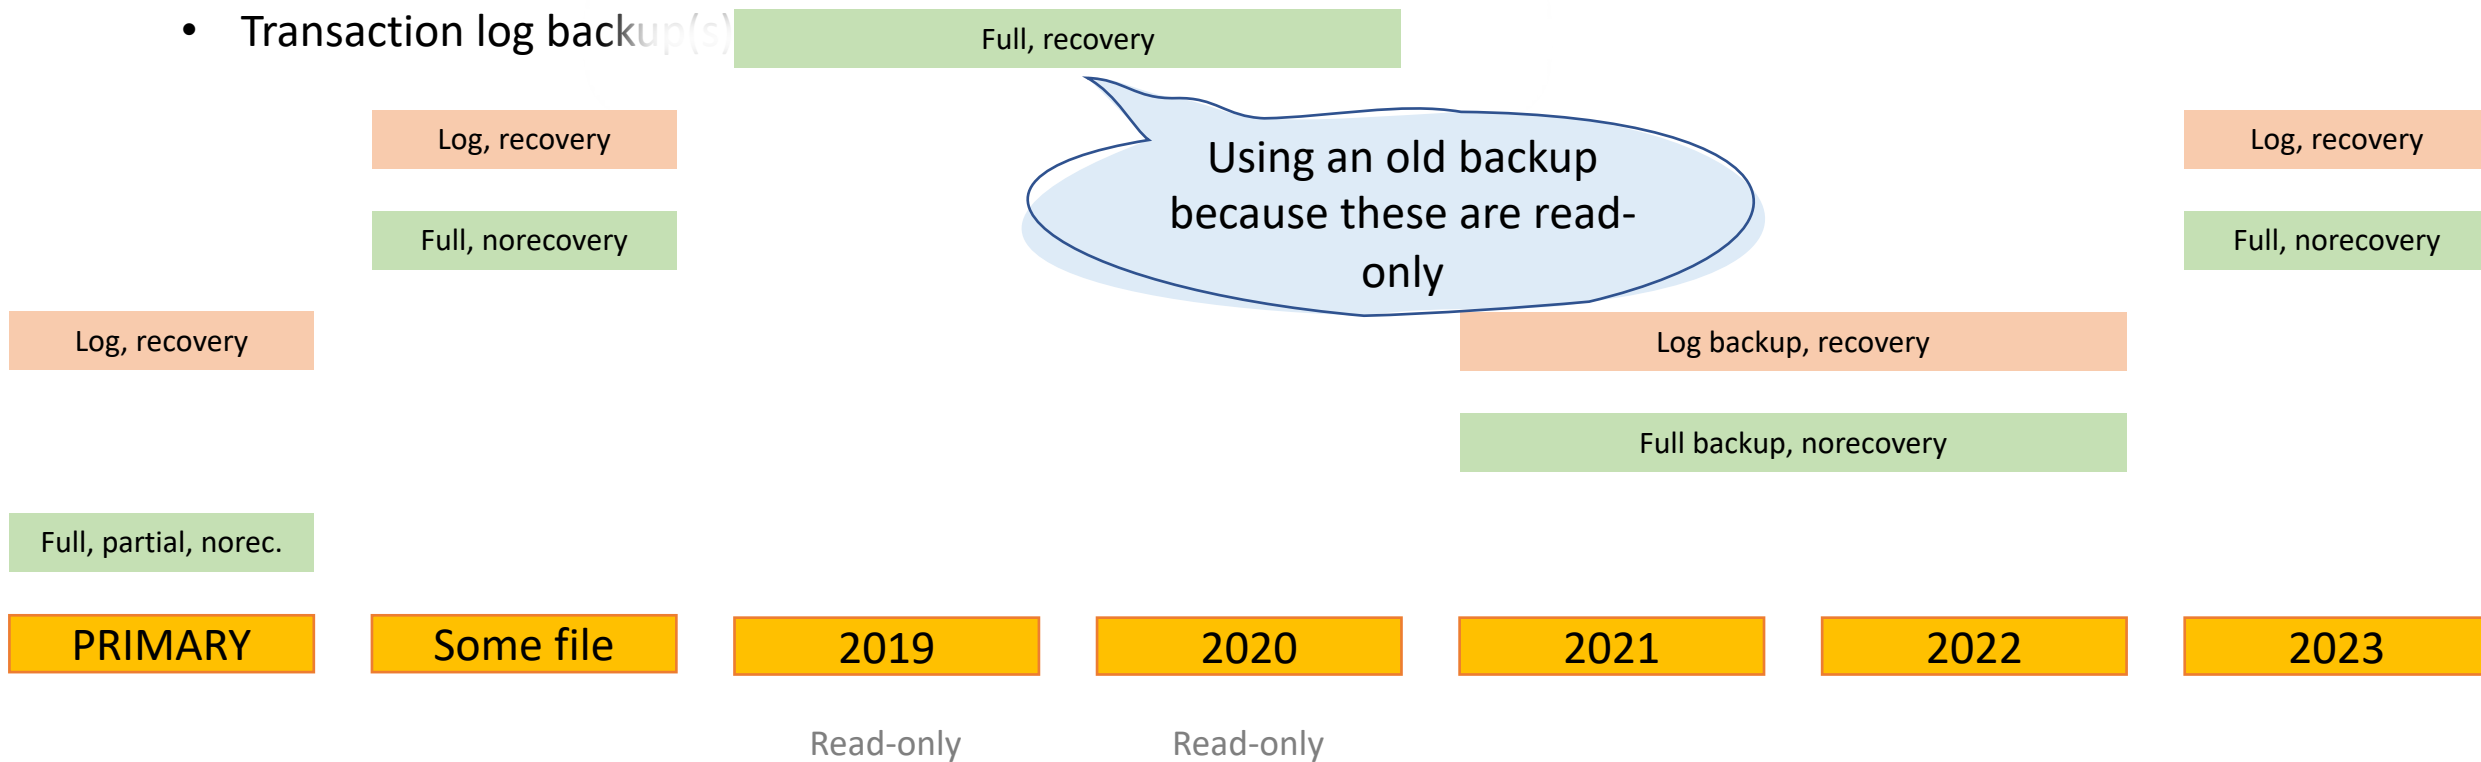
- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups
- Transaction log backup(s) for the read-write filegroups



# Piecemeal restore (full + log)

We have:

- An old FULL backup of the read-only filegroups
- A recent FULL backup of the read-write filegroups
- Transaction log backup(s)







DEMO



# Seeks & scans work differently

```
SELECT AccountID, SUM(Amount) AS Net_balance_change  
FROM dbo.AccountTransactions  
WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
      AND AccountID BETWEEN 810000000 AND 820000000  
GROUP BY AccountID;
```

# Seeks & scans work differently

```
SELECT AccountID, SUM(Amount) AS Net_balance_change  
FROM dbo.AccountTransactions  
WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
      AND AccountID BETWEEN 810000000 AND 820000000  
GROUP BY AccountID;
```

AccountID
TransactionDate
TransactionID

# Seeks & scans work differently

```
SELECT AccountID, SUM(Amount) AS Net_balance_change
FROM dbo.AccountTransactions
WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'
      AND AccountID BETWEEN 810000000 AND 820000000
GROUP BY AccountID;
```

AccountID
TransactionDate
TransactionID

YEAR(TransactionDate)
AccountID
TransactionDate
TransactionID



# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

AccountID
TransactionDate
TransactionID

YEAR(TransactionDate)
AccountID
TransactionDate
TransactionID

29 600 pages

- Seek AccountID 810000000-820000000

# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

AccountID
TransactionDate
TransactionID

YEAR(TransactionDate)
AccountID
TransactionDate
TransactionID

29 600 pages

- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2022-06

# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

AccountID
TransactionDate
TransactionID

29 600 pages

YEAR(TransactionDate)
AccountID
TransactionDate
TransactionID

- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2022-06
- Aggregate on AccountID

# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

AccountID
TransactionDate
TransactionID

29 600 pages

- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2022-06
- Aggregate on AccountID

YEAR(TransactionDate)
AccountID
TransactionDate
TransactionID

1 800 pages

- For partition 2021
- Seek AccountID 810000000-820000000

# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

AccountID
TransactionDate
TransactionID

29 600 pages

- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2022-06
- Aggregate on AccountID

YEAR(TransactionDate)
AccountID
TransactionDate
TransactionID

1 800 pages

- For partition 2021
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2021-12

# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

AccountID
TransactionDate
TransactionID

29 600 pages

- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2022-06
- Aggregate on AccountID

YEAR(TransactionDate)
AccountID
TransactionDate
TransactionID

1 800 pages

- For partition 2021
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2021-12
- Aggregate on AccountID

# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

AccountID
TransactionDate
TransactionID

29 600 pages

- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2022-06
- Aggregate on AccountID

YEAR(TransactionDate)
AccountID
TransactionDate
TransactionID

1 800 pages

- For partition 2021
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2021-12
- Aggregate on AccountID

600 pages

- For partition 2022
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2022-01 – 2022-06
- Aggregate on AccountID

# Seeks & scans work differently

```
WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
    AND AccountID BETWEEN 810000000 AND 820000000
```

- Seek AccountID 810000000-820000000

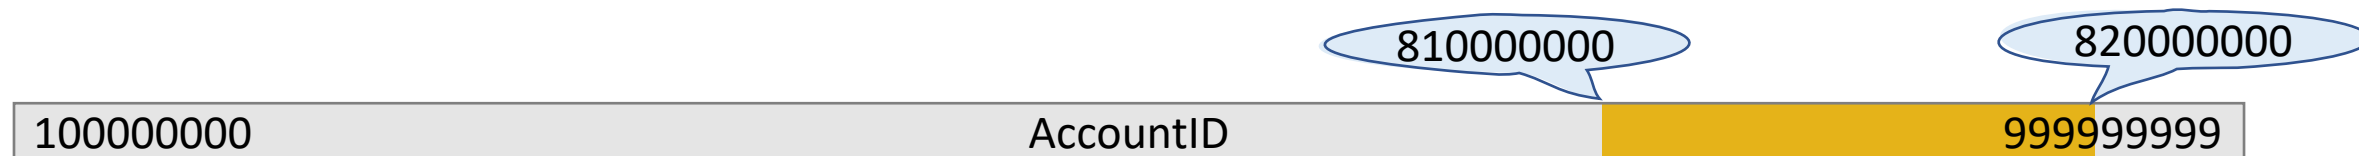
100000000	AccountID	999999999
-----------	-----------	-----------



# Seeks & scans work differently

```
WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
    AND AccountID BETWEEN 810000000 AND 820000000
```

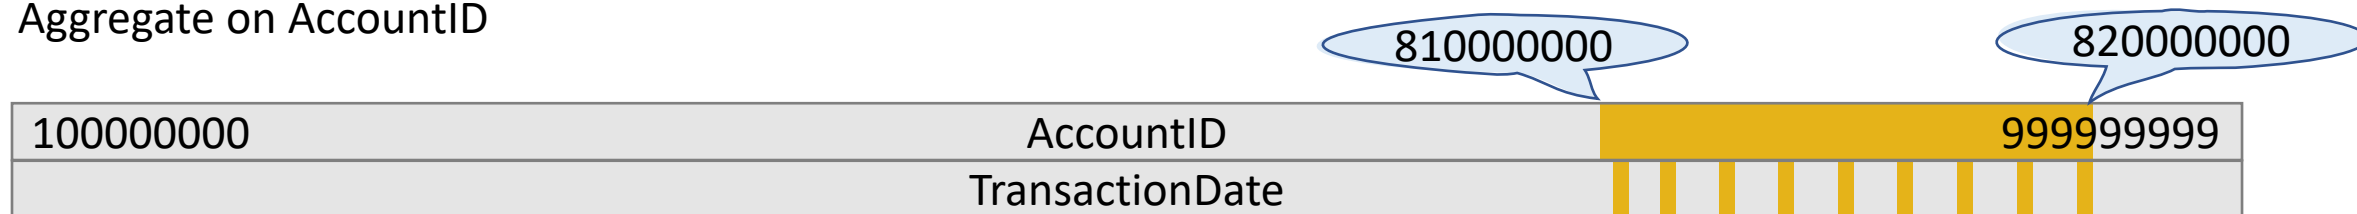
- Seek AccountID 810000000-820000000



# Seeks & scans work differently

```
WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000
```

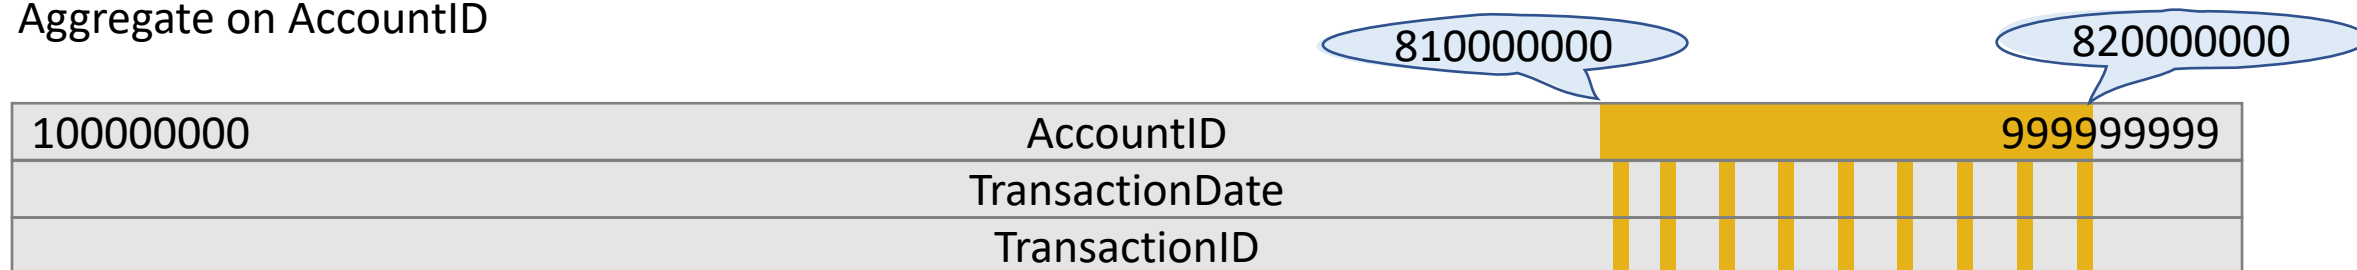
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2022-06
- Aggregate on AccountID



# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

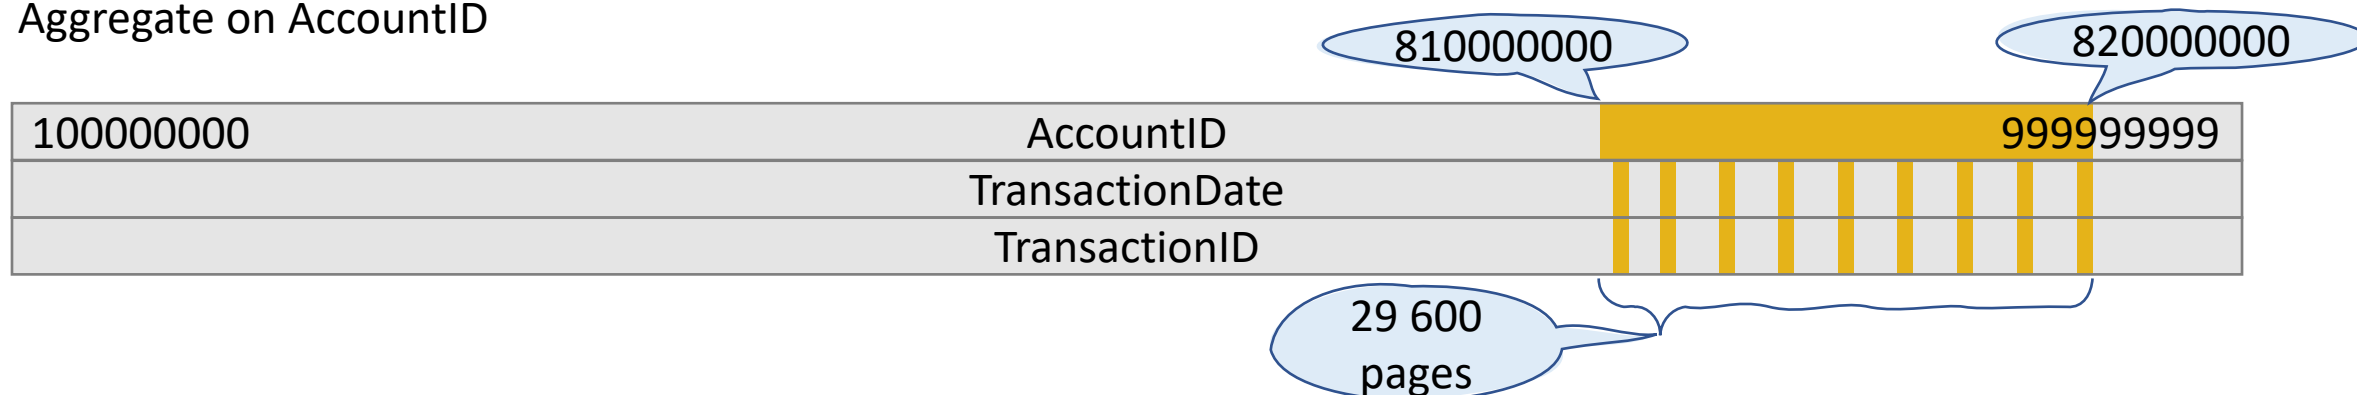
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2022-06
- Aggregate on AccountID



# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

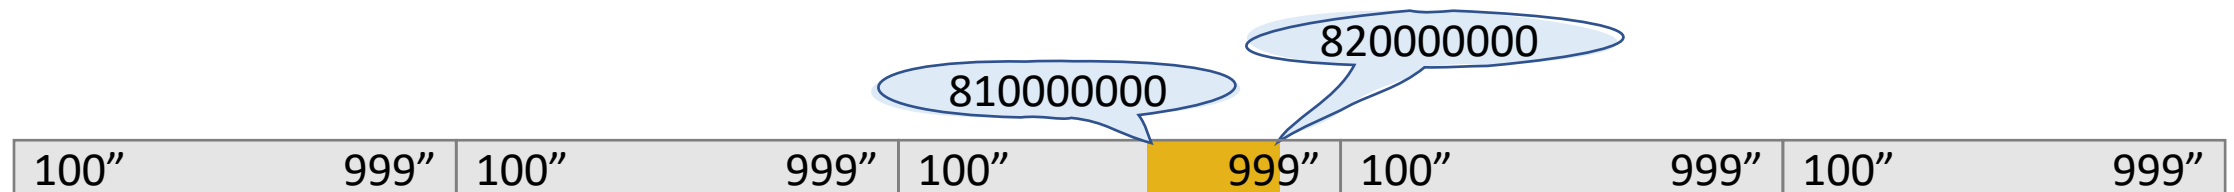
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2022-06
- Aggregate on AccountID



# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

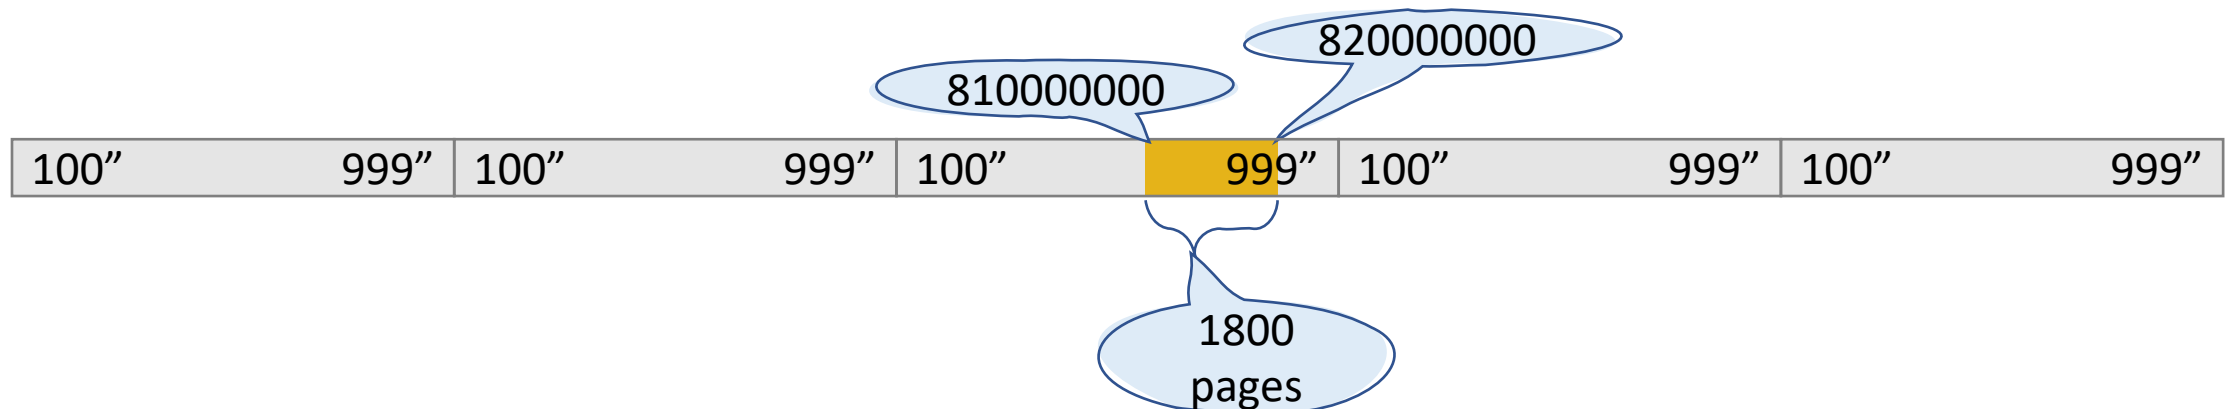
- For partition 2021
- Seek AccountID 810000000-820000000



# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

- For partition 2021
- Seek AccountID 810000000-820000000



## Seeks & scans work differently

```
WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'
      AND AccountID BETWEEN 810000000 AND 820000000
```

- For partition 2021
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2021-12
- Aggregate on AccountID

The diagram shows a table with 10 columns. The first row contains values: 100'', 999'', 100'', 999'', 100'', 999'', 100'', 999'', 100'', 999''. The second row is highlighted in yellow and contains the text TransactionDate. The third row is also highlighted in yellow and contains the text TransactionID. Two callouts point to the 6th and 7th columns of the first row, with values 810000000 and 820000000 respectively.

100''	999''	100''	999''	100''	999''	100''	999''	100''	999''
TransactionDate									
TransactionID									

# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

- For partition 2021
  - Seek AccountID 810000000-820000000
  - Filter TransactionDate 2021-07 – 2021-12
  - Aggregate on AccountID
- For partition 2022
  - Seek AccountID 810000000-820000000

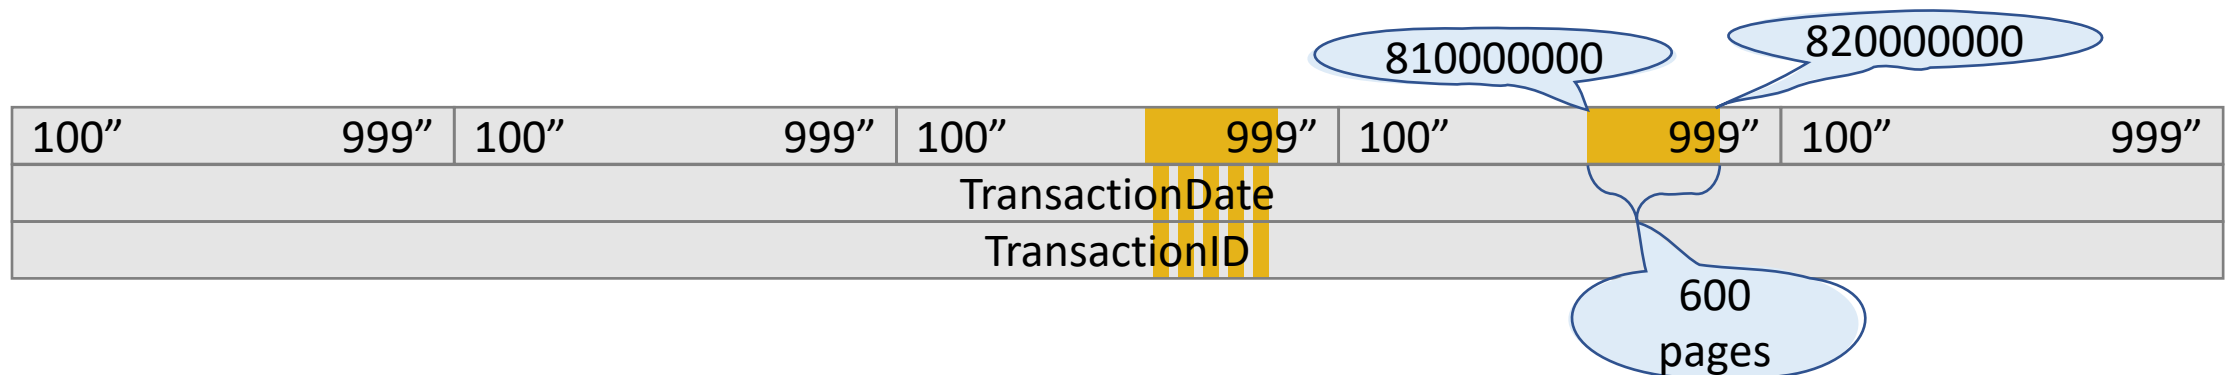
100"	999"	100"	999"	100"	999"	100"	999"	100"	999"
TransactionDate									
TransactionID									



# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

- For partition 2021
  - Seek AccountID 810000000-820000000
  - Filter TransactionDate 2021-07 – 2021-12
  - Aggregate on AccountID
- For partition 2022
  - Seek AccountID 810000000-820000000

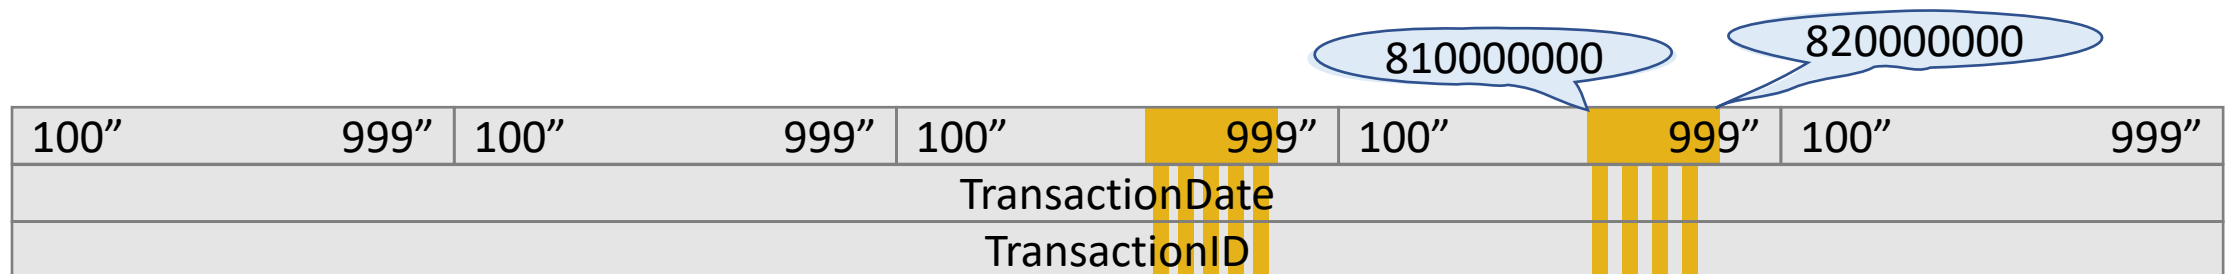


# Seeks & scans work differently

WHERE TransactionDate BETWEEN '2021-07-01' AND '2022-06-30'  
AND AccountID BETWEEN 810000000 AND 820000000

- For partition 2021
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2021-12
- Aggregate on AccountID

- For partition 2022
- Seek AccountID 810000000-820000000
- Filter TransactionDate 2021-07 – 2021-12
- Aggregate on AccountID



# DEMO



# Takeaway: Why partitioning?

- Some costly DML operations can be turned into super-fast DDL statements
- Some queries can be tuned to use a single partition (partition elimination)
- Can solve specific latch contention issues by spreading writes
- You can backup only read-write partitions – faster backups
- You can restore the important partitions first – shorter RTO

... or maybe not?

- More complex to maintain
- More complex to develop
- Tuned for non-partitioned  $\neq$  tuned for partitioning

# Standing on the shoulders of giants

- Kendra Little
- Cathrine Wilhelmsen
- Uwe Ricken

# One more thing



1. Go to [evals.datagrillen.com](https://evals.datagrillen.com)
2. Find my session in the list
3. Write something nice

Email: [daniel@strd.co](mailto:daniel@strd.co)  
Twitter: [@dhmacher](https://twitter.com/dhmacher)  
Blog: [sqlsunday.com](https://sqlsunday.com)